



# The Future of AI-Driven Software Engineers

AI Master Class



# The Future of AI-Driven Software Engineers



## Thomas Martinsen

Tech Evangelist @ Twoday + Microsoft Regional Director & Microsoft AI MVP



## Frederik Baun Hansen

Global Senior Architect @ Twoday



## Henri Schulte

Partner Solution Architect @ Microsoft



## Andreas Hoffmann

Software Developer @ Twoday

# THANK YOU TO OUR AMAZING SPONSORS !

glueck■kanja

robopack   
empowered by SOFTWARE  
CENTRAL

nerdio

ARROW

Ontinue

veeam

 VARONIS

eido

 Admin By Request  
ZERO TRUST PLATFORM

Lenovo

 dynatrace

 VirtualMetric

 Red Hat

Arkimentum

 semperis

DanofficeIT

 EPIC FUSION  
BRING IT ALL TOGETHER

Recast

 twoday

 European  
Power Platform  
Conference  
COPENHAGEN 29 JUNE - 02 JULY 2024

 inforcer

 APENTO

 TRUESEC

 GLOBE SYSTEMS

 controlUP

 PATCH  
MY PC

JN DATA  
WE ARE ON

 systemcenterdukes

 EasyEntra

 Microsoft

 Implora.io



 TD SYNNEX

 TOKEN  
token2.swiss

 Cloud Factory

 GLOBETEAM

 NORDSTERN

 VENZO\_  
a valantic company

 FEITIAN  
WE BUILD SECURITY

 identity stack

 2LINKIT

# Guidelines for this master class

## Be active and deliberate

Engage directly in the exercises and discussions rather than observing from the sidelines. The value of this master class comes from applying the concepts in practice, challenging assumptions, and testing ideas in real time. Approach the case work with intent: define goals clearly, articulate constraints, and avoid jumping to solutions before the problem is well understood.

## Work with structure

Treat ideas, outputs, and decisions with an engineering mindset. Be explicit about assumptions, boundaries, and trade-offs. Approach problems methodically, and validate results instead of accepting them at face value. Use structure in how you think, collaborate, and document your work. Clear reasoning and disciplined execution are central to making AI-driven development reliable.

## Connect to your reality

Continuously reflect on how the methods and patterns discussed apply to your own organization, systems, and responsibilities. Leverage the different perspectives in your team, especially across development and infrastructure roles. Aim to leave with concrete actions you can implement immediately, as well as a clearer understanding of what must change to adopt these practices effectively.

# Takeaways for this master class

## Engineering shifts from coding to intent and specs

Modern AI-native development moves the engineer's focus from writing code to defining intent, specifications, architecture boundaries, and validation gates. The spec becomes the core unit of work, enabling predictable and controlled agent execution.

## Reliable AI development requires structure, not ad-hoc prompts

Professional AI-assisted development depends on structured prompt engineering, agent primitives, and context engineering. Treating prompts and agent behavior as reusable, versioned engineering assets makes outcomes predictable and scalable.

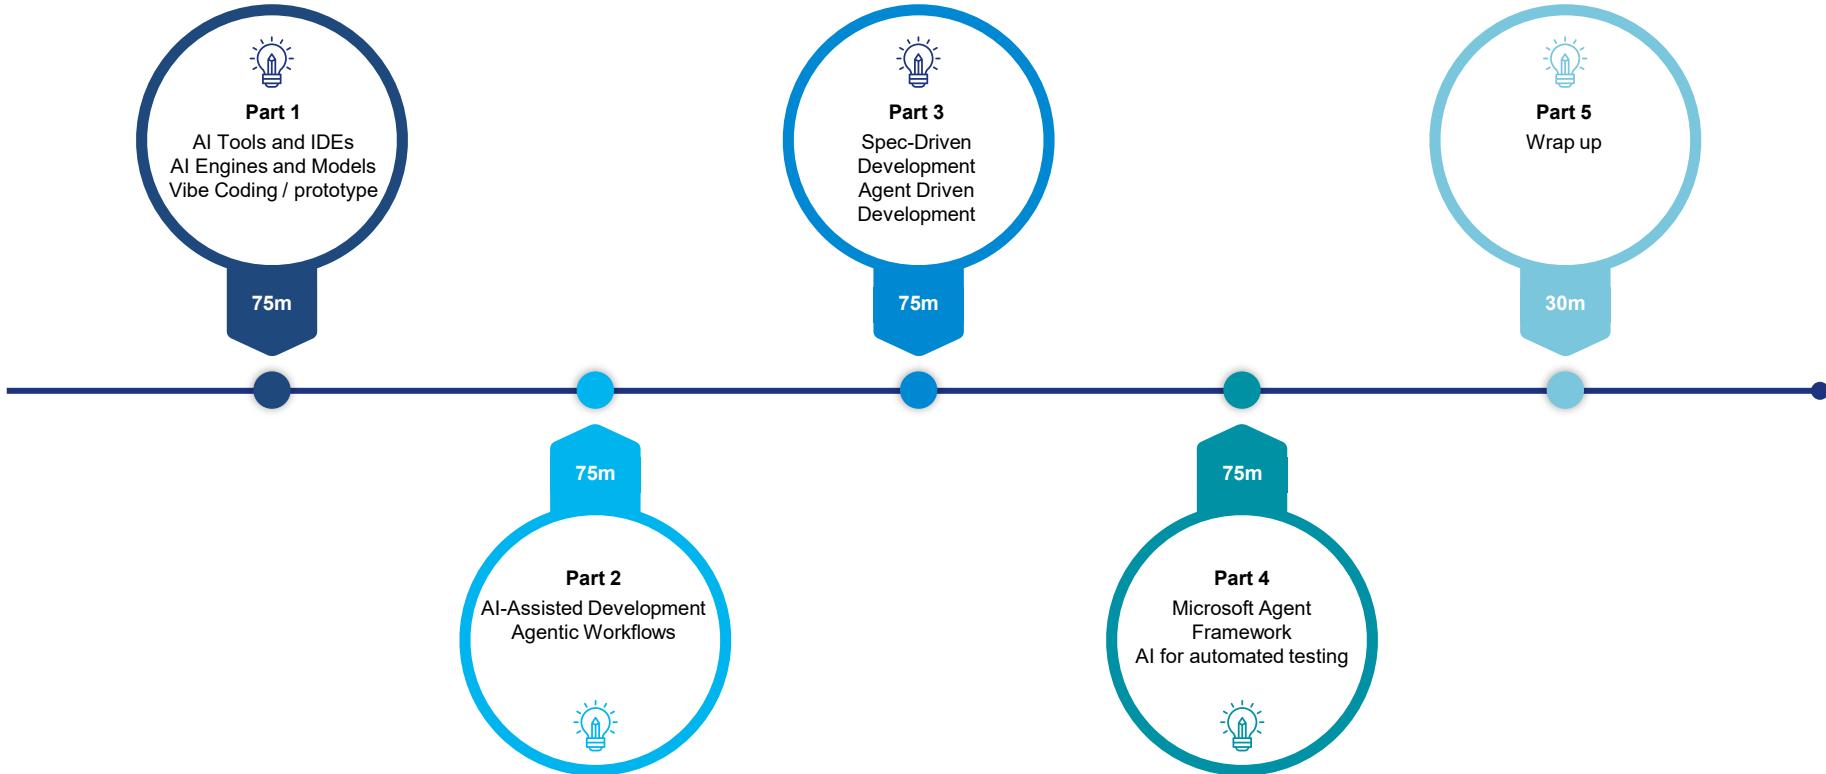
## AI supports the full software lifecycle—not just code

AI can do more than generate code: it can drive automated testing, documentation, validation, and other lifecycle tasks. These capabilities can be orchestrated end-to-end using interoperable frameworks like the Microsoft Agent Framework, with human oversight and validation gates.

# GitHub Repos

<https://github.com/ELDK26-AI-masterclass>

# Agenda





# Custom Tees, Powered by AI

Shop unique, smartly designed T-shirts.

[Shop Now](#)



## Treadley Case Overview

### AI Tools and Models

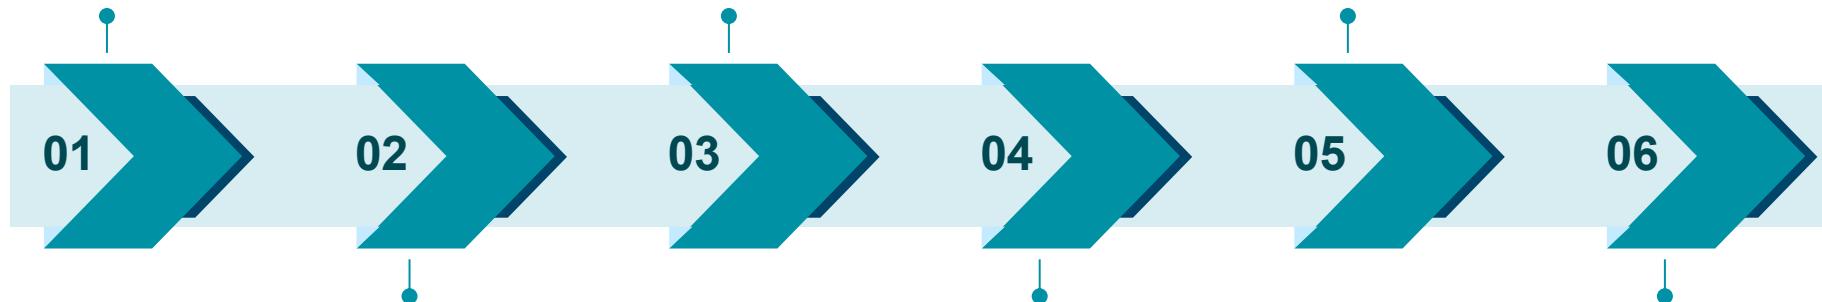
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

### AI-Assisted Development

- Prompt engineering beyond "do X"

### Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

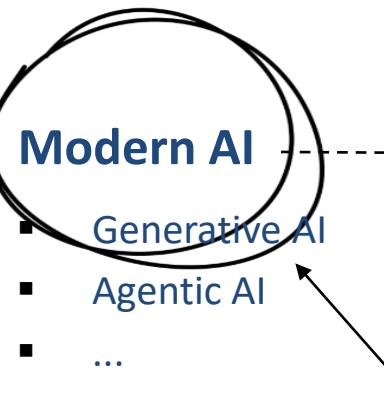
- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior

# Introduction to modern software development

# Modern AI

## Traditional AI

- Rule-based Systems
- Symbolic AI
- Machine Learning
- Deep Learning
- Regression
- Classification
- Forecasting



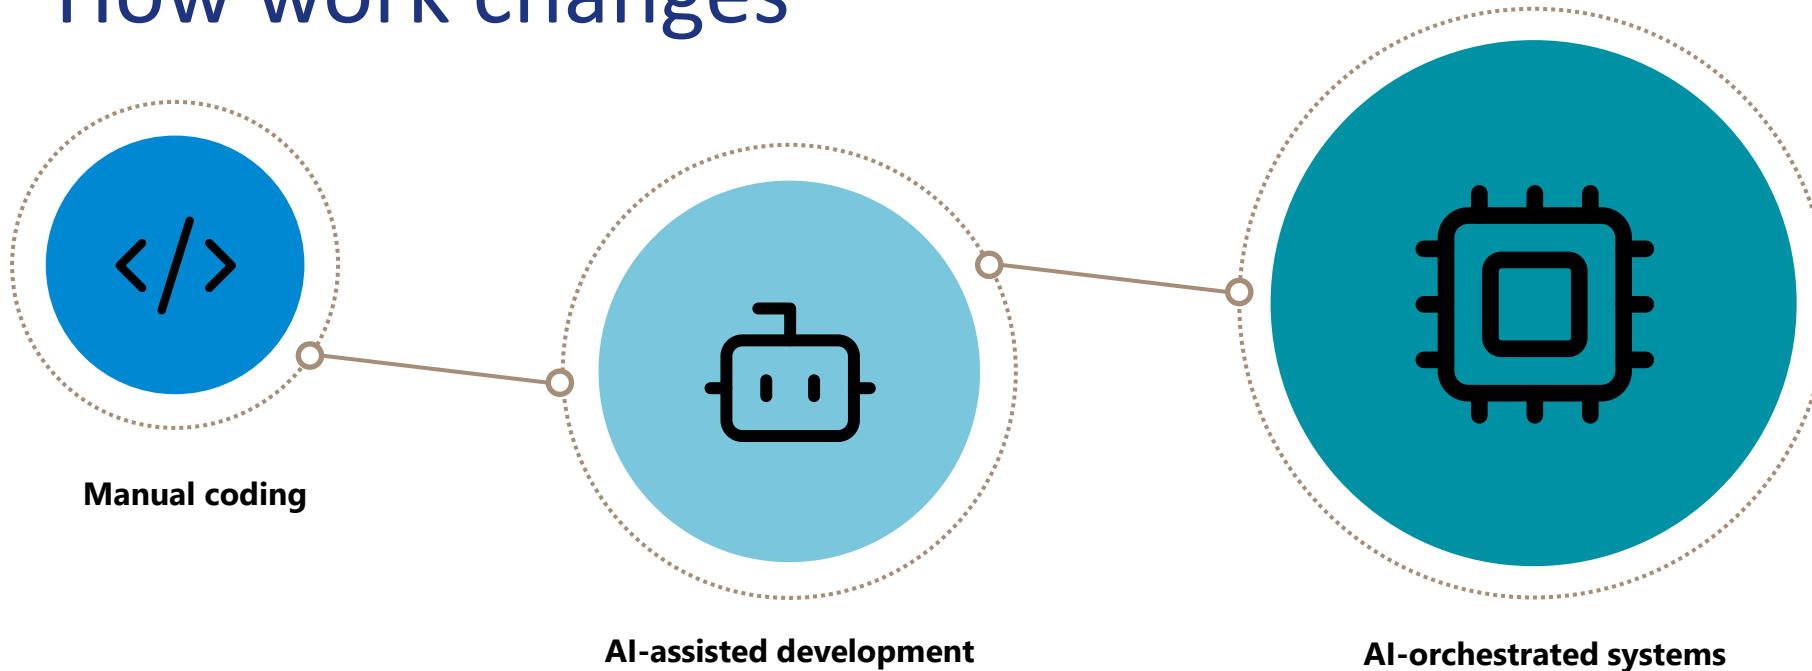
## Strong AI

- AGI
- Super Intelligence
- Singularity

Relevant for all companies  
and all employees – and all  
software engineers



# How work changes



- Developer writes everything
- Manual scaffolding
- Hand-written tests
- Large code surface

Producer → Intent designer + reviewer →

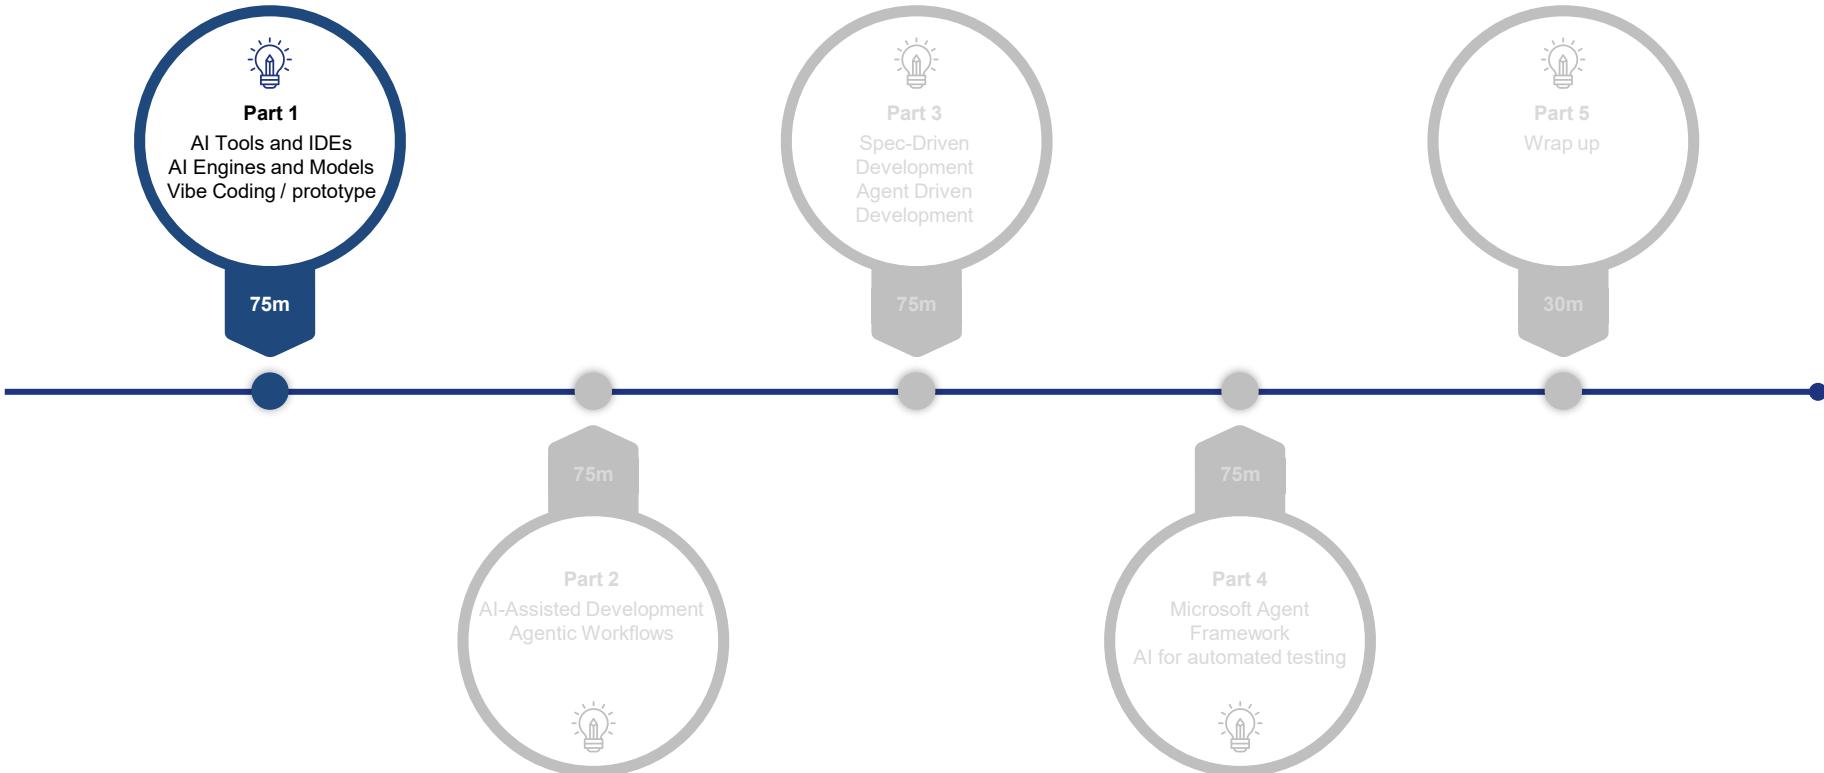
- Developer defines intent
- Spec-driven planning
- Agent execution
- Human validation gates



CASE

# Getting started

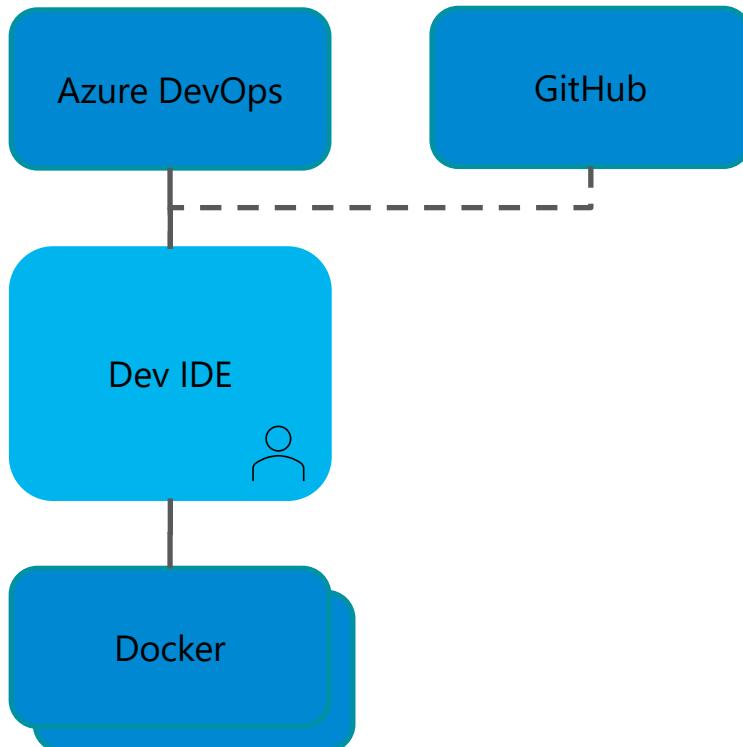
# Agenda



# AI Tools and IDEs

# AI Engines and Models

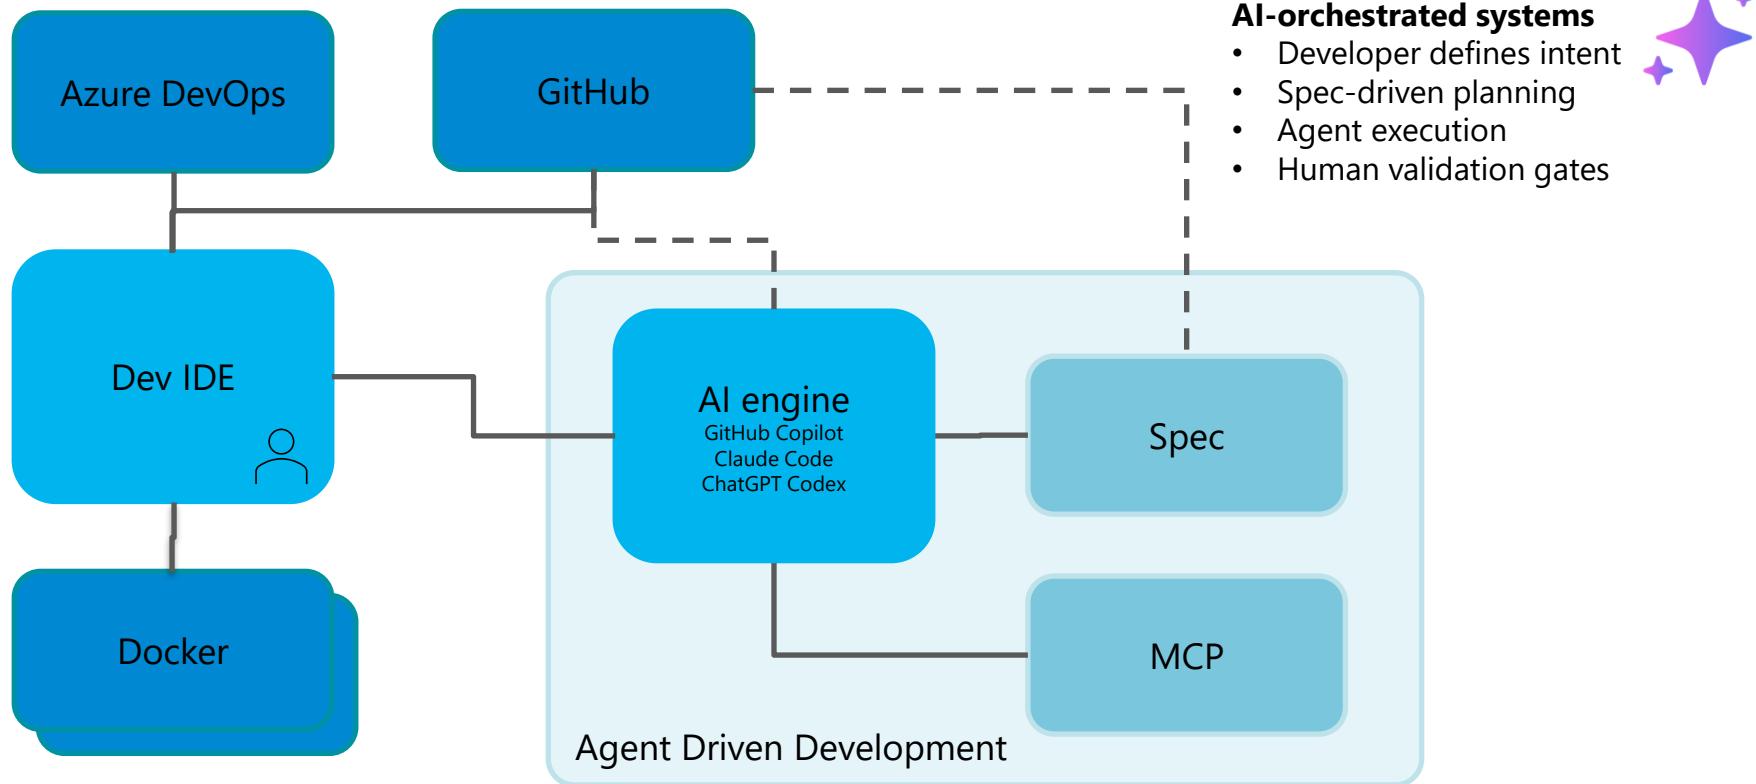
# Manual coding setup



## Manual coding

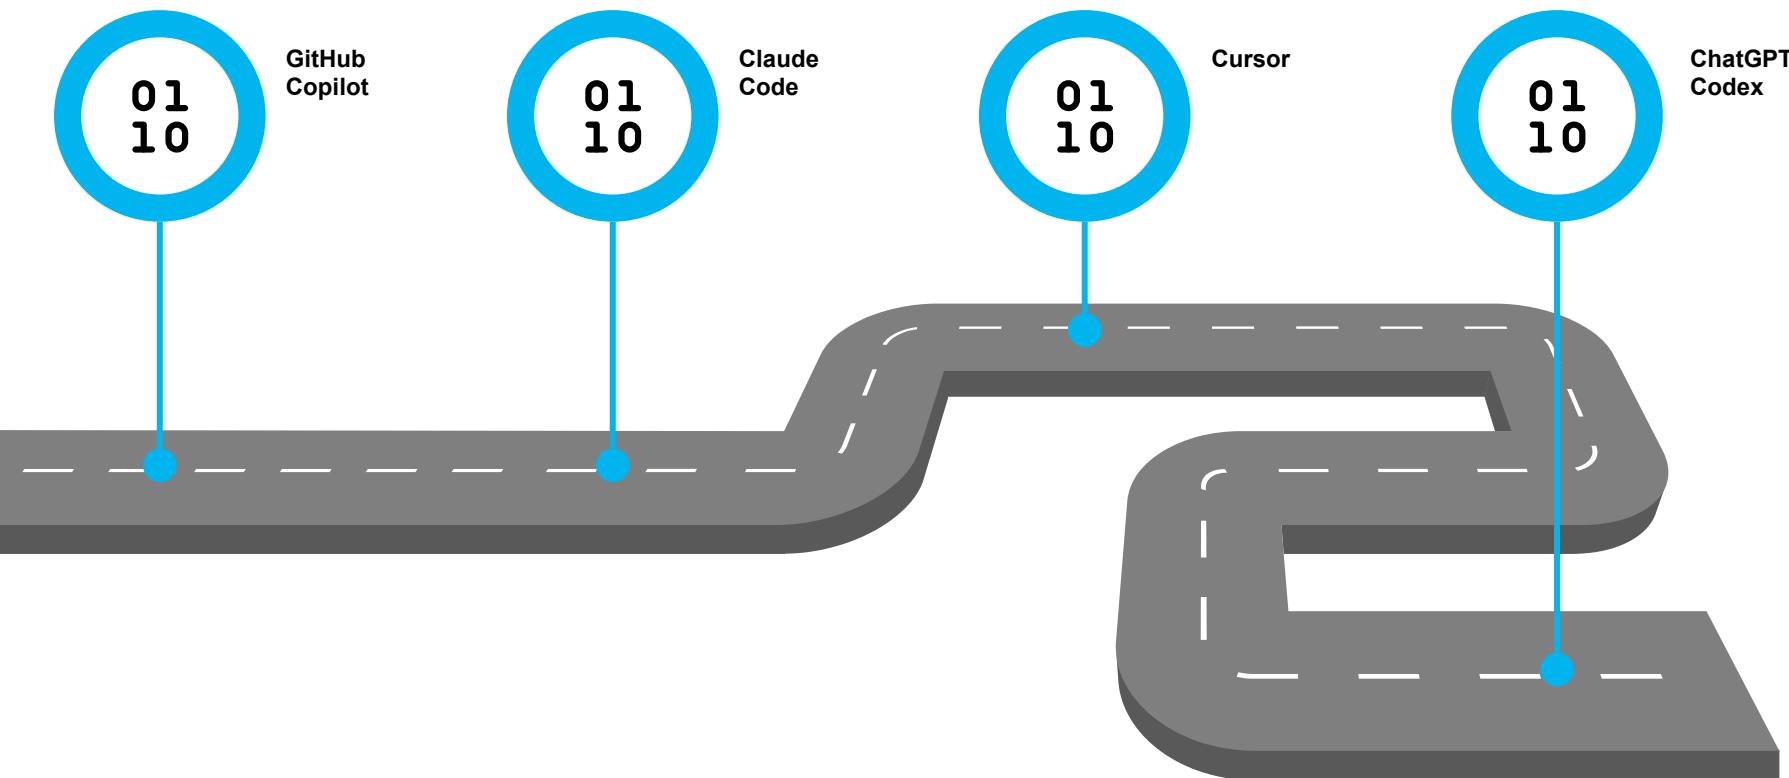
- Developer writes everything
- Manual scaffolding
- Hand-written tests
- Large code surface

# AI-orchestration setup

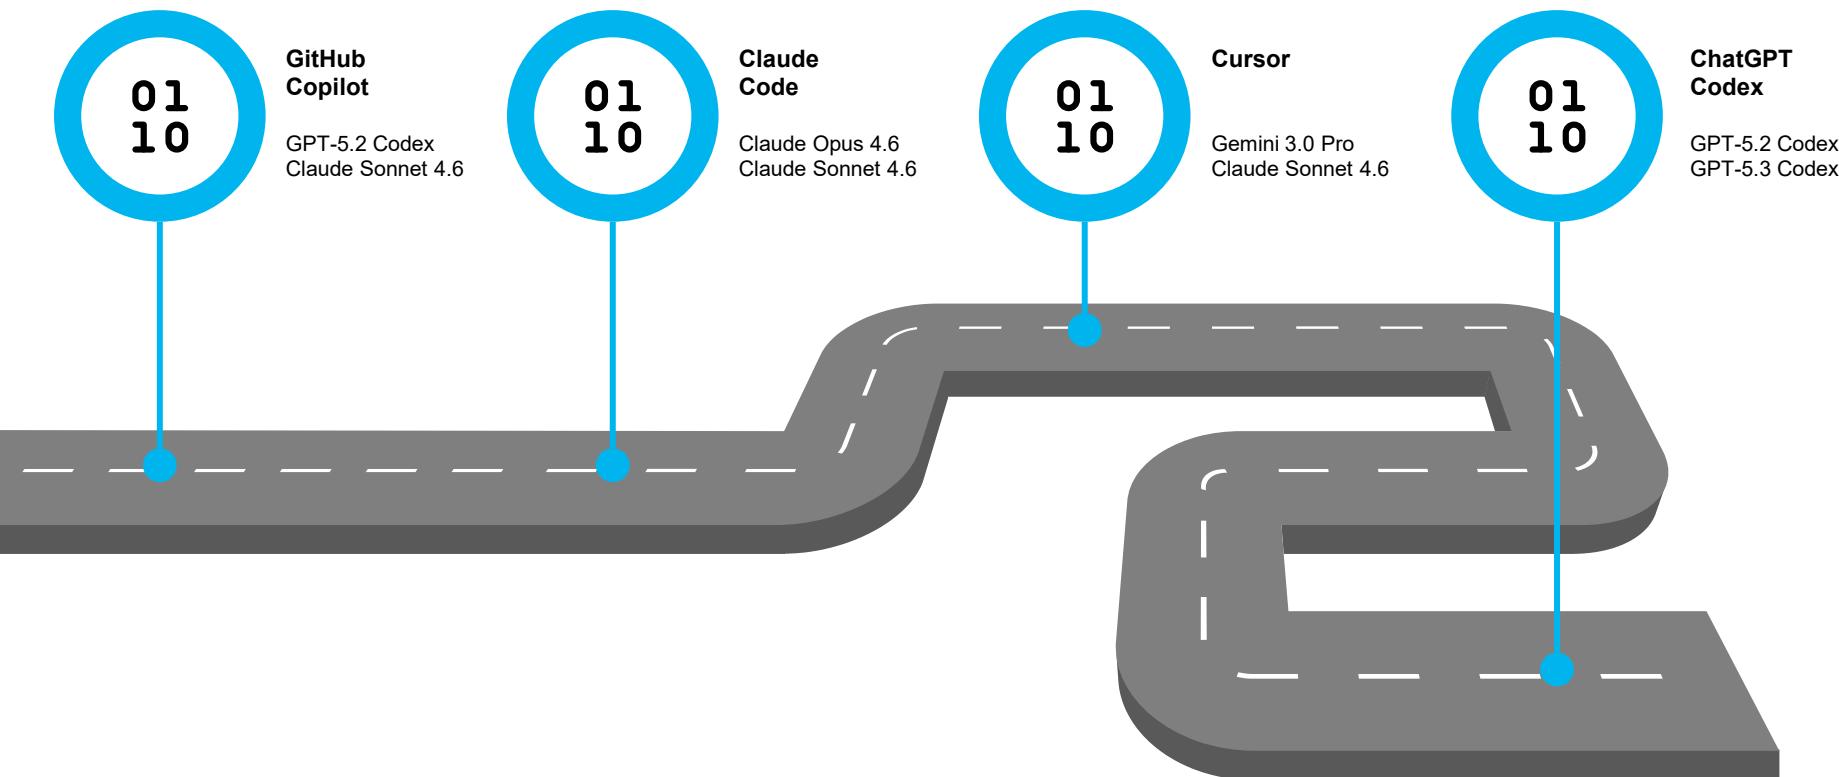




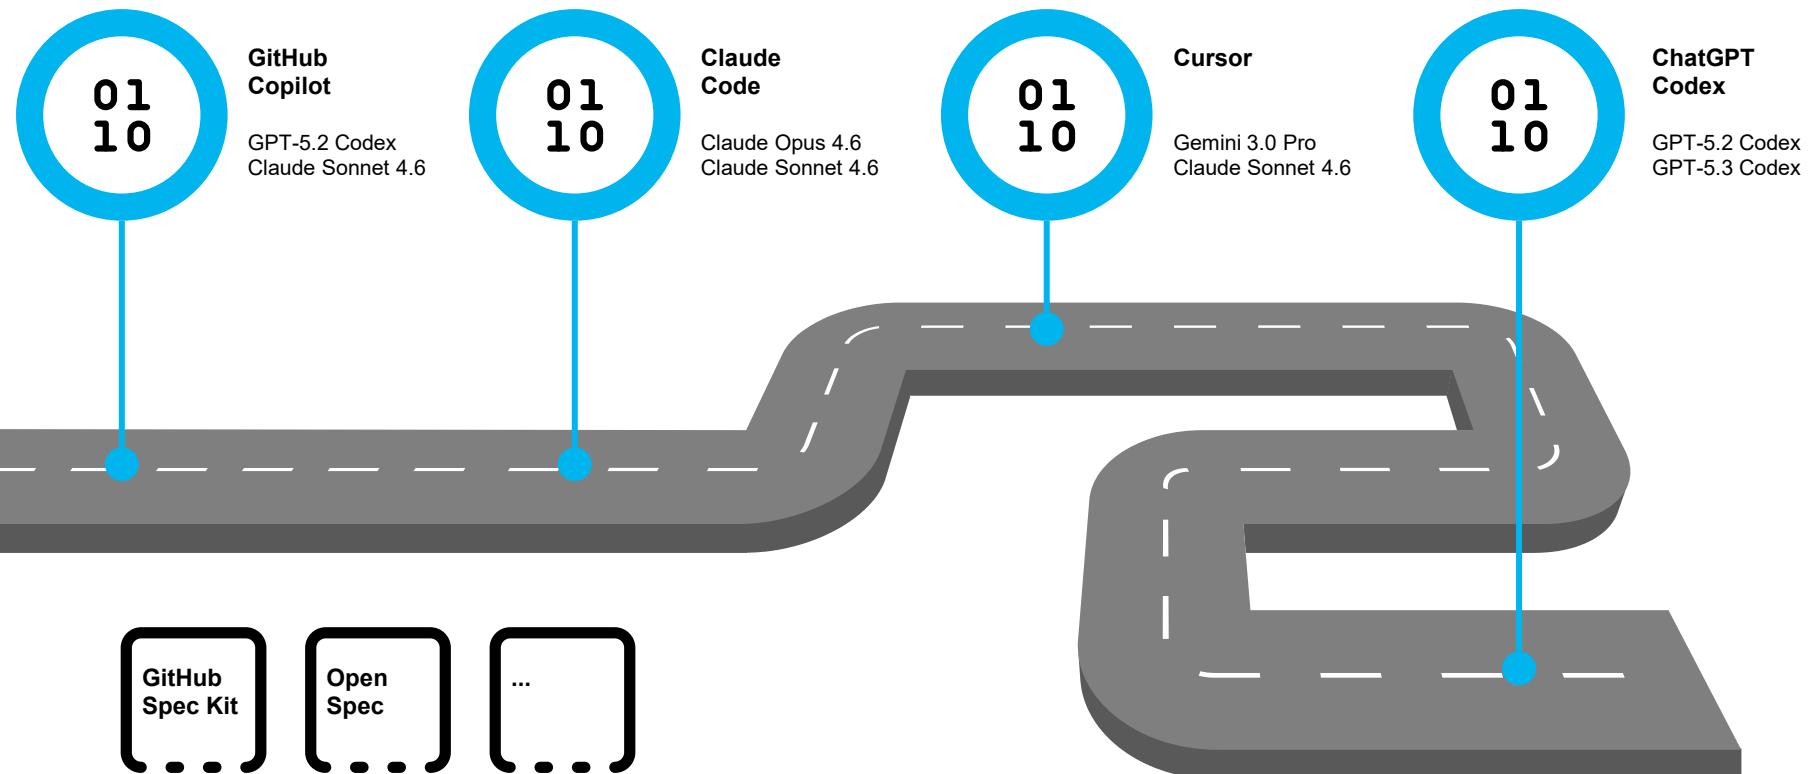
# Know your AI tools



# Know your AI tools and your models

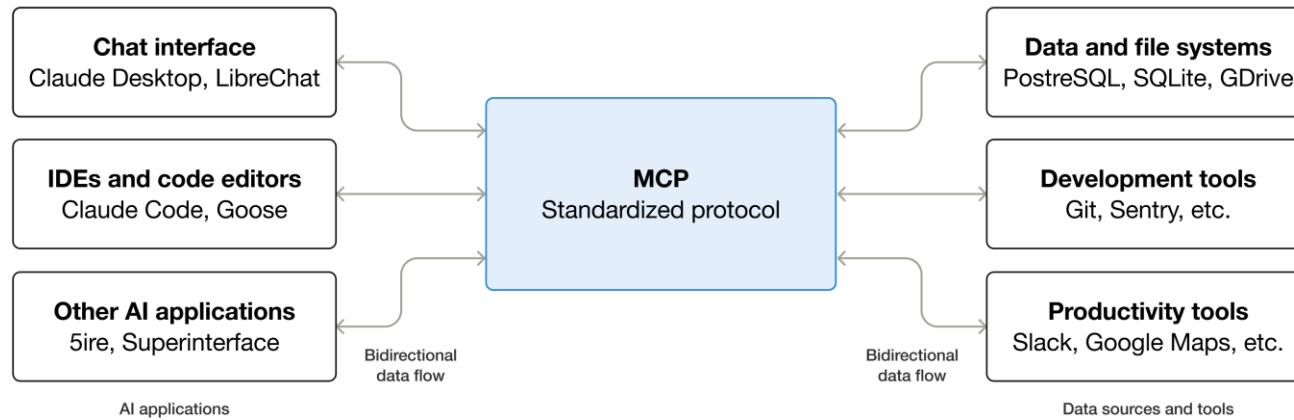


# Know your AI tools and your models





# MCP





D E M O

# AI tools and models in action

GitHub Copilot + ChatGPT Codex + OpenCode + OpenSpec + MCP



CASE

AI tools and models



# AI Tools and Models

## AI Tools and Models

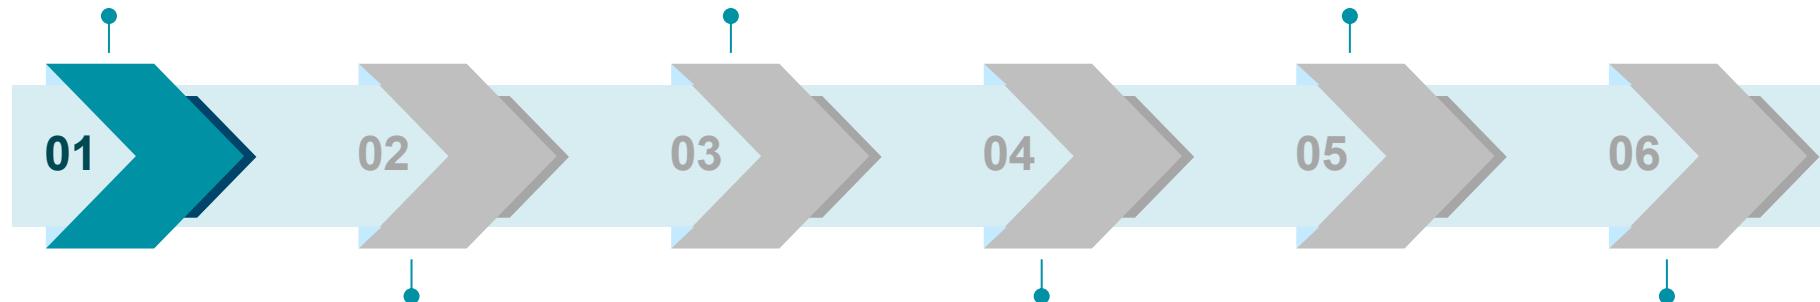
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

## AI-Assisted Development

- Prompt engineering beyond "do X"

## Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



### Vibe Coding

- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

### Agentic Workflows

- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

### Agent Framework

- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior



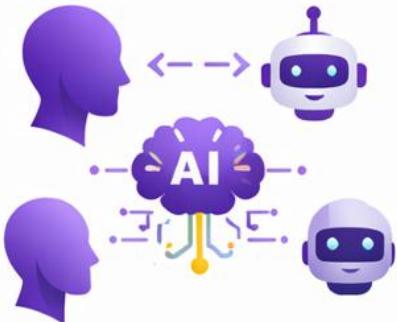
# Vibe Coding

# Why Vibe Coding exists

- Software complexity keeps increasing
- Software systems are growing in complexity
- AI is good at repetition and wiring
- Developers spend significant time on mechanical work

Human  
Intent  
Judgment

AI  
Generation  
Iteration



*Vibe coding shifts focus from  
typing code to **describing** outcomes*



# What is Vibe Coding?

- Intent-driven development
- Conversational, iterative process
- Output quality depends on human input

Vibe Coding = describing **what** you want,  
how it should **feel** and what it should **look** like



*"The quality of the result depends on the **quality** of the direction"*





D E M O

Vibe Coding with GitHub Spark

# Vibe Coding Platforms Today

Level 1 – Inside the IDE

- Cursor
- Windsurf

Level 2 – Solution-Level builders

- GitHub Spark

Level 3 – Product-Level Generators

- Figma Make
- Lovable
- Emergent



*Same mindset – different abstraction levels*

# What Vibe Coding is – and is not

It is:

- Fast feedback
- Great for exploration
- A new way to collaborate
- Accelerating prototype and MVP solutions

It is not:

- Thinking-free
- Guaranteed clean architecture
- Free of responsibility



*"Because something is **simple** – doesn't necessarily make it **easy**"*



CASE

## Vibe Coding

<https://github.com/ELDK26-AI-masterclass/case>



# Vibe Coding

## AI Tools and Models

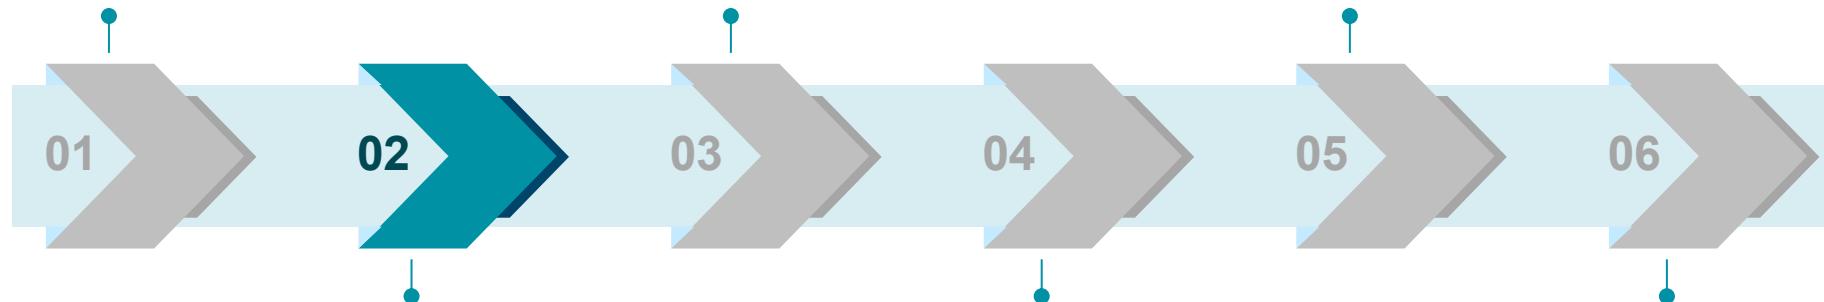
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

## AI-Assisted Development

- Prompt engineering beyond "do X"

## Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



### Vibe Coding

- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

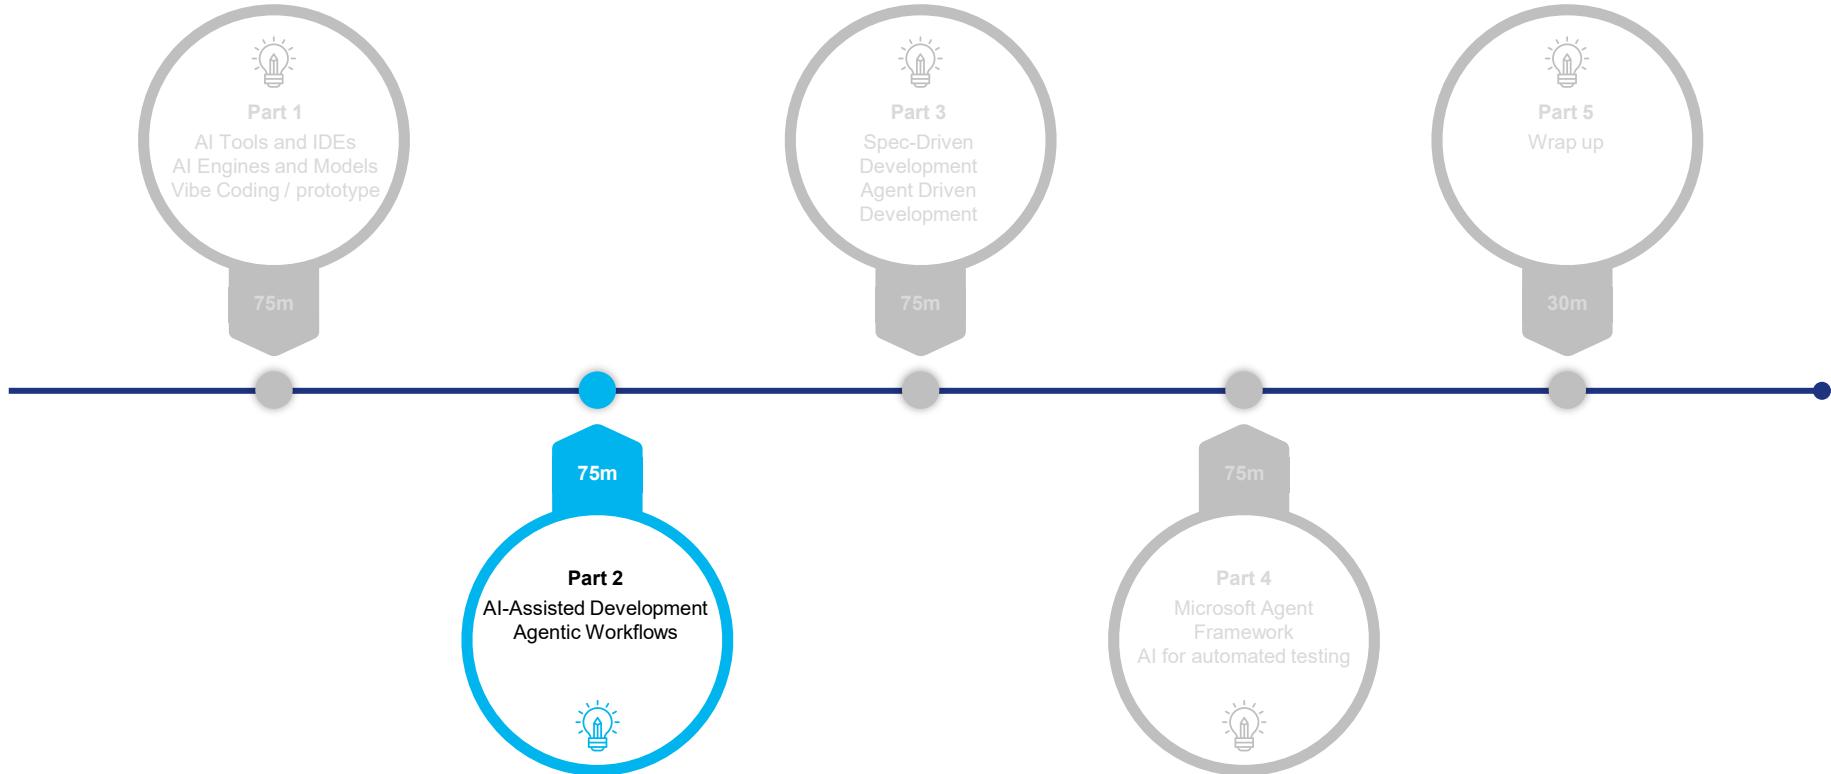
### Agentic Workflows

- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

### Agent Framework

- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior

# Agenda

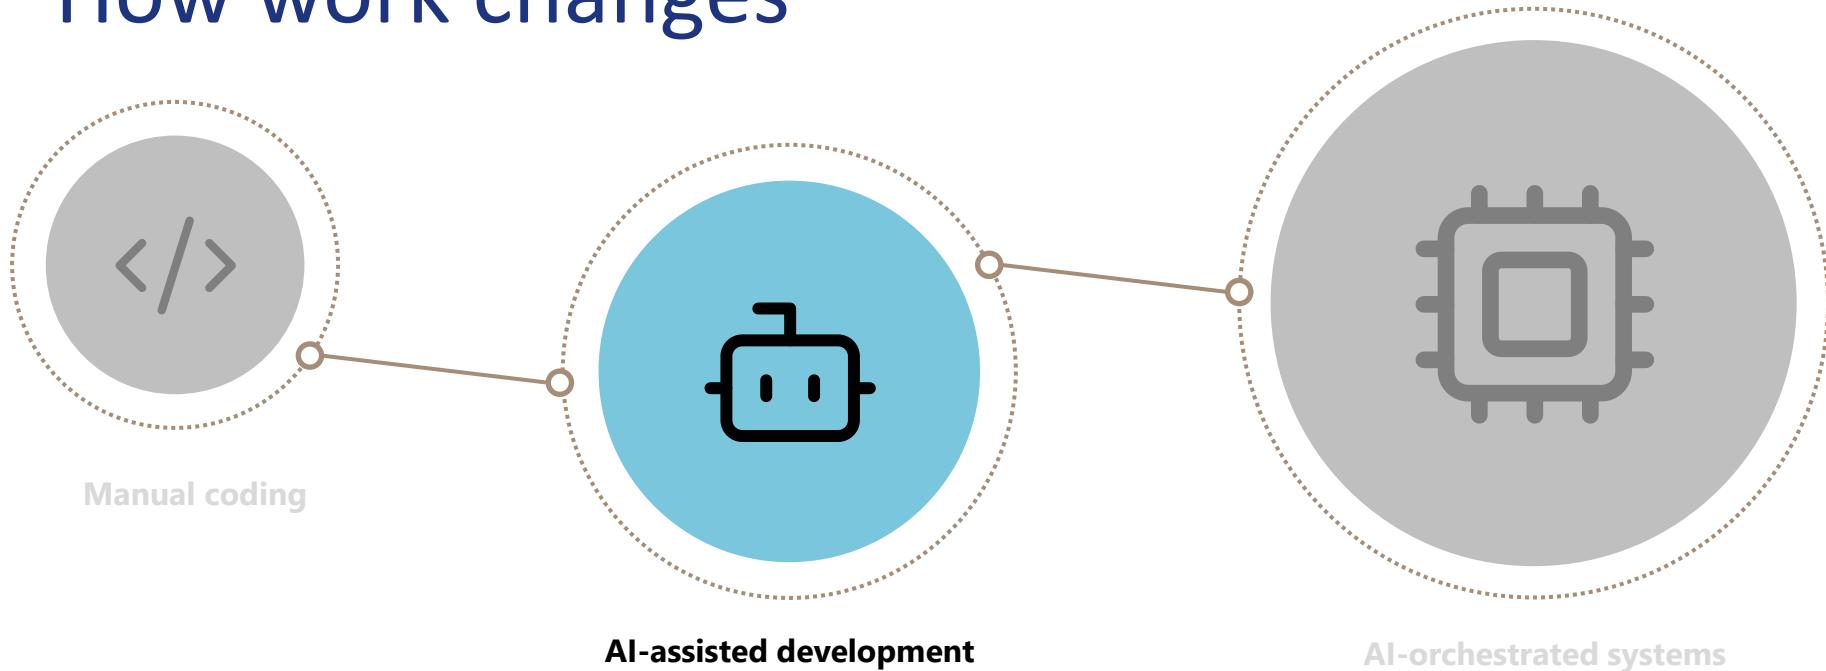




# AI-assisted development



# How work changes



- Developer writes everything
- Manual scaffolding
- Hand-written tests
- Large code surface

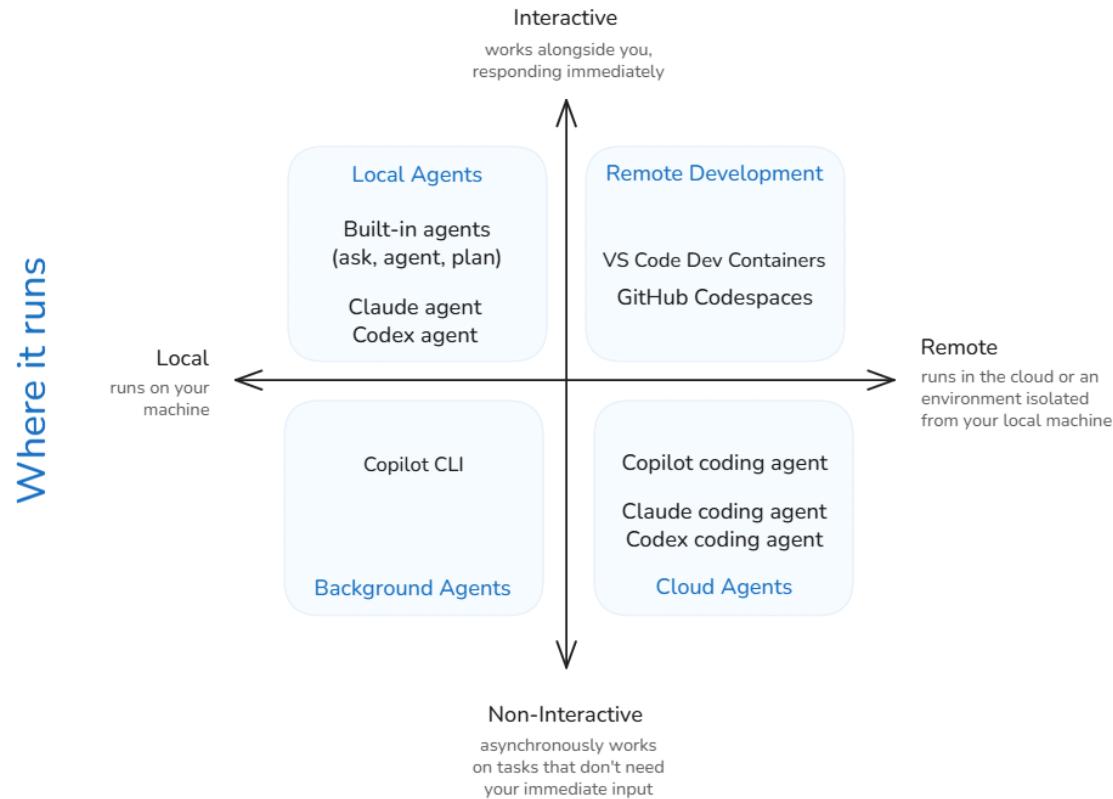
Producer → Intent designer + reviewer →

- Developer defines intent
- Spec-driven planning
- Agent execution
- Human validation gates



# Agent Sessions

## How you collaborate





D E M O

# Agent Sessions

Local + Remote



# How agents run

## Local

### Local execution in the IDE

- Best when: early iteration
- Benefits: fast feedback
- Example: building first .spec.md

## Async

### Async delegation

- Best when: repetitive tasks
- Benefits: runs in background
- Example: test generation

## Hybrid

### Hybrid orchestration

- Best when: multi-agent flows
- Benefits: scalable implementation
- Example: feature spec → tests → docs



CASE

AI-assisted development



# AI-Assisted Development

## AI Tools and Models

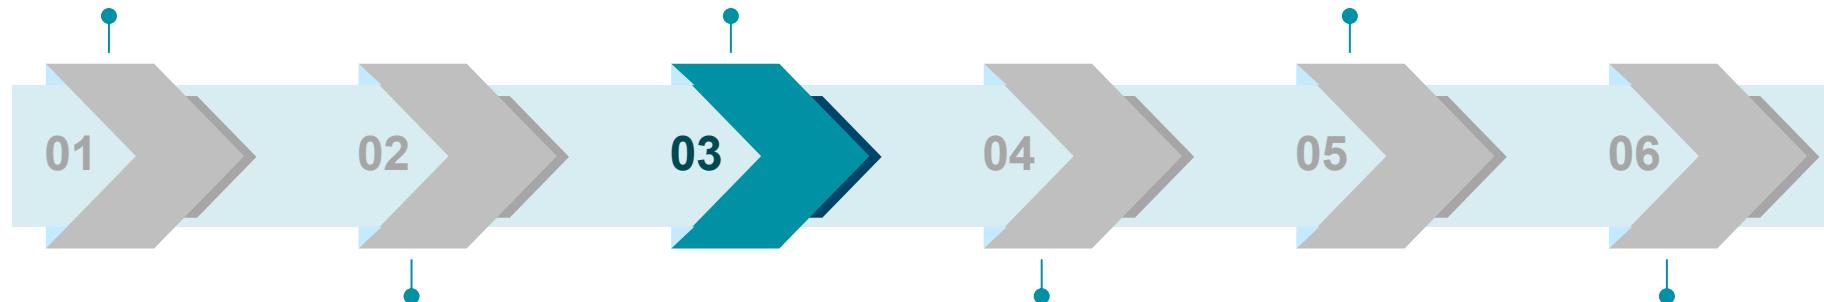
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

## AI-Assisted Development

- Prompt engineering beyond "do X"

## Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



### Vibe Coding

- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

### Agentic Workflows

- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

### Agent Framework

- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior



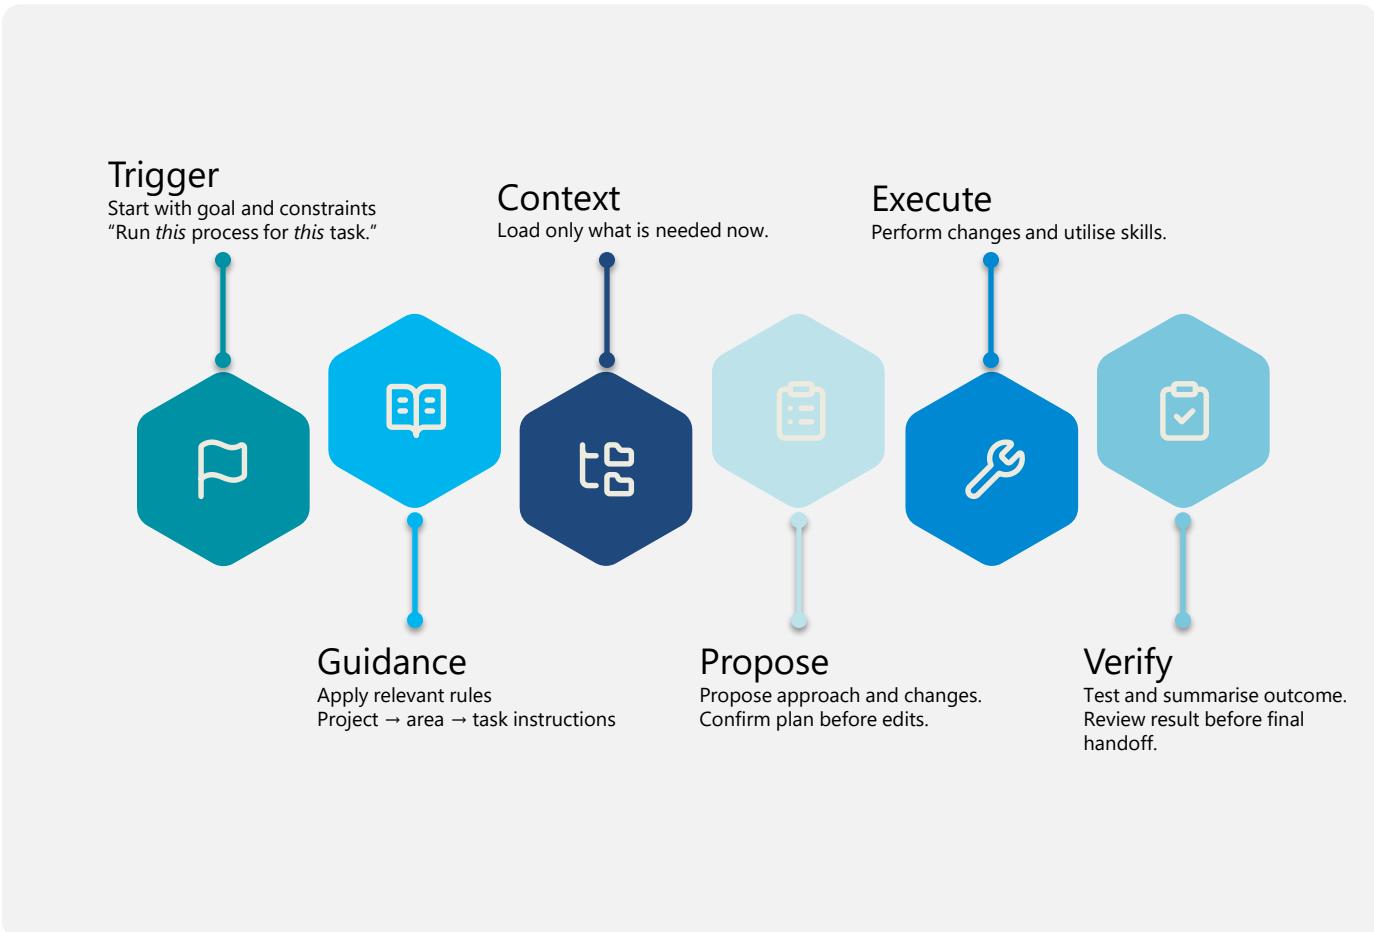
# Agentic Workflows



# Agentic Workflows

Agentic workflow is a reusable process encoded as a prompt.

- Captured in a workflow prompt
- Orchestrates Agent Primitives
- Runs in phases
- Goal is consistent outcomes

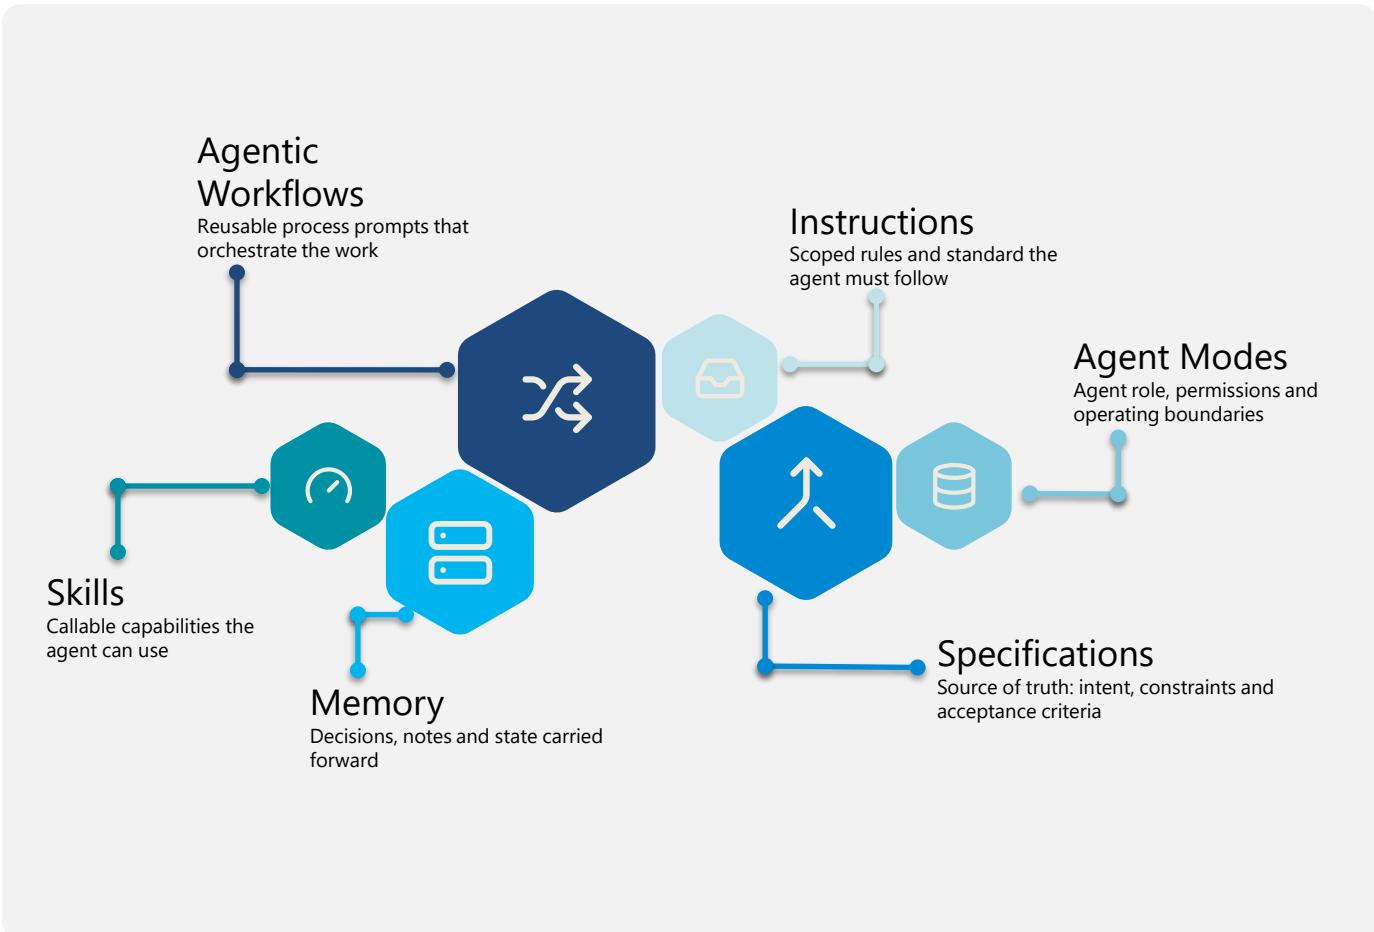




# Agentic Building Blocks

Agentic workflows scale when we externalize behavior into project assets.

Agentic Building Blocks make work repeatable, reviewable and shared across the team.





D E M O

# Agentic workflows

Instructions + Specifications + Skills



CASE

Agentic workflows



# Agentic Workflows

## AI Tools and Models

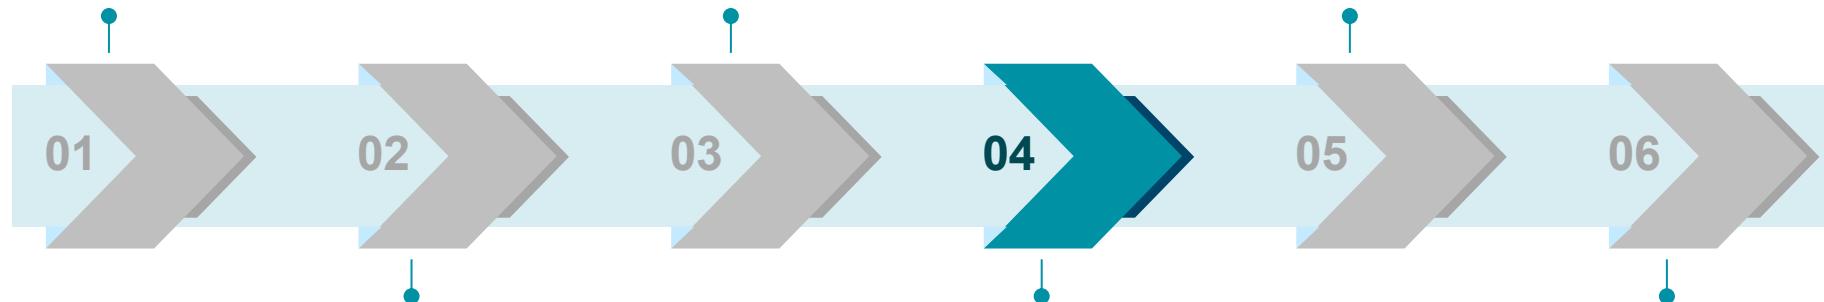
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

## AI-Assisted Development

- Prompt engineering beyond "do X"

## Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



### Vibe Coding

- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

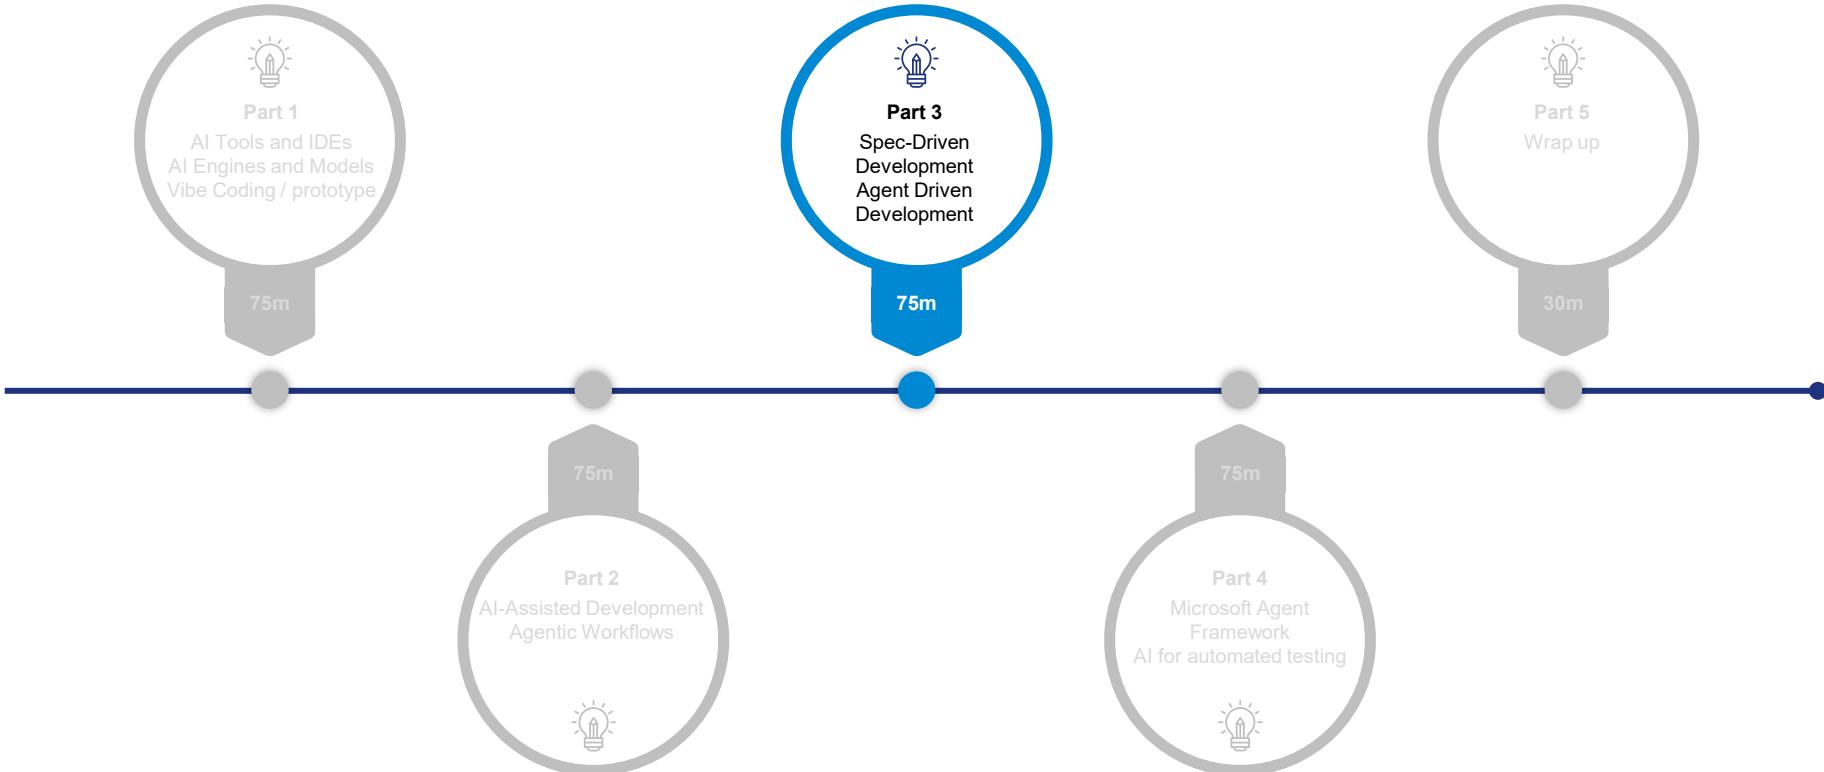
### Agentic Workflows

- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

### Agent Framework

- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior

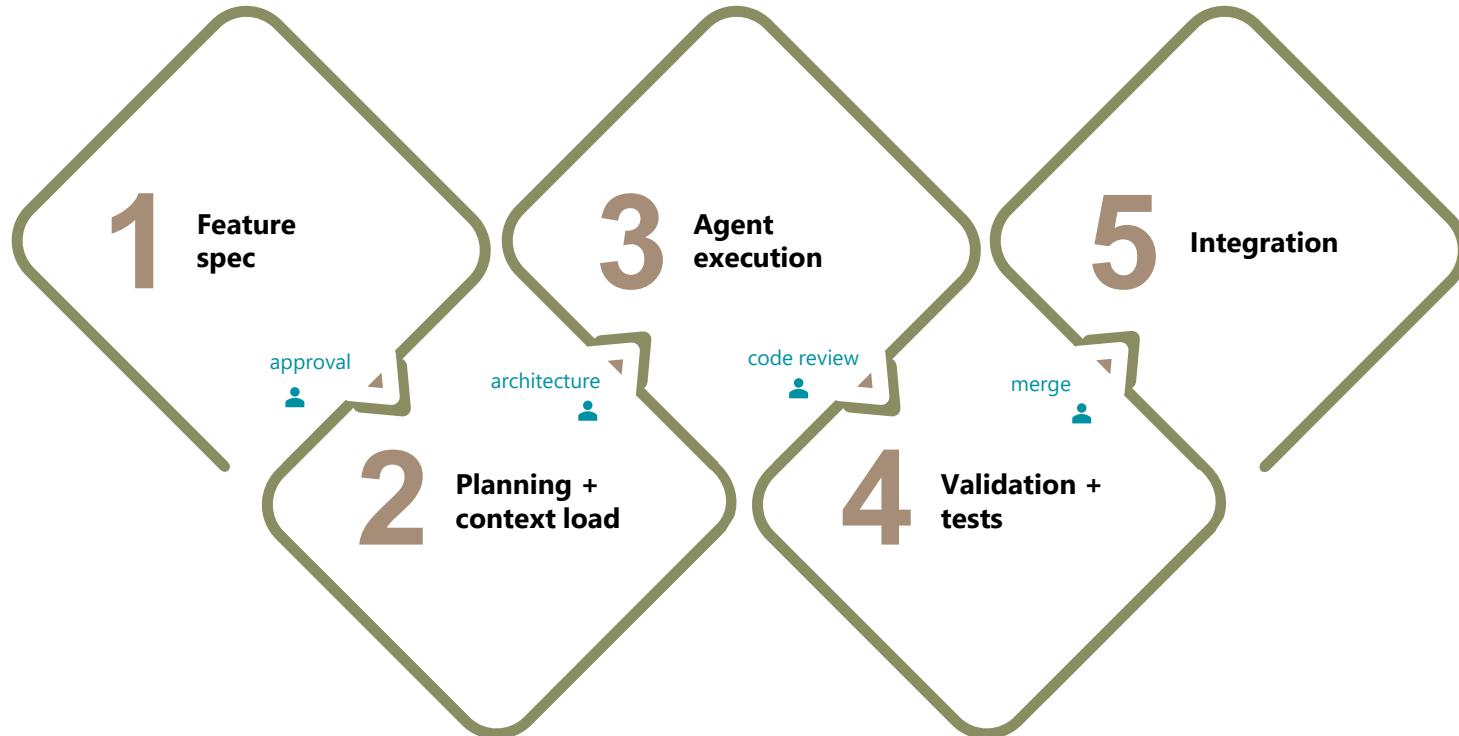
# Agenda



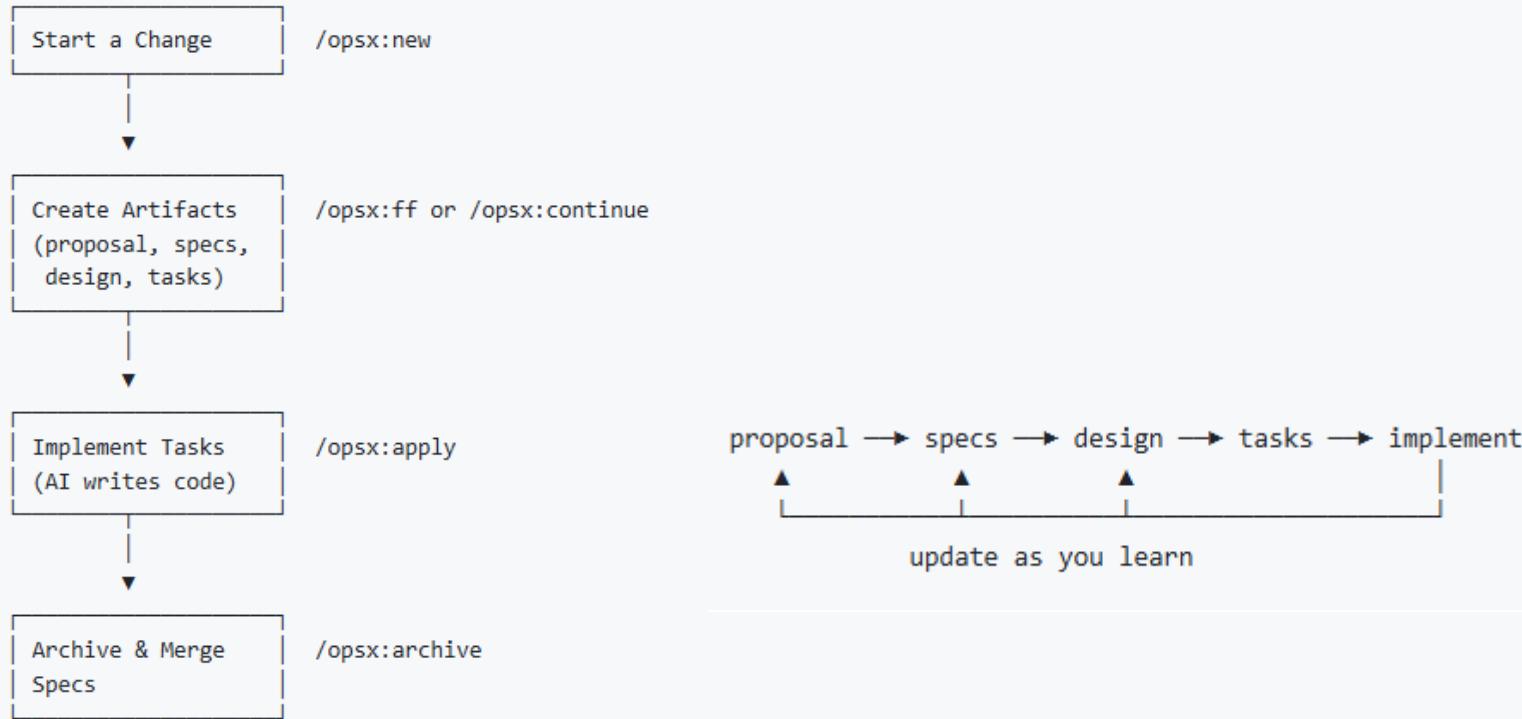
# Spec-Driven Development



# Spec-driven development



# OpenSpec



# Change – Implement – Archive

```
You: /opsx:new add-dark-mode
```

```
AI: Created openspec/changes/add-dark-mode/  
Ready to create: proposal
```

```
You: /opsx:ff # "fast-forward" - generate all planning docs
```

```
AI: ✓ proposal.md - why we're doing this, what's changing  
✓ specs/ - requirements and scenarios  
✓ design.md - technical approach  
✓ tasks.md - implementation checklist  
Ready for implementation!
```

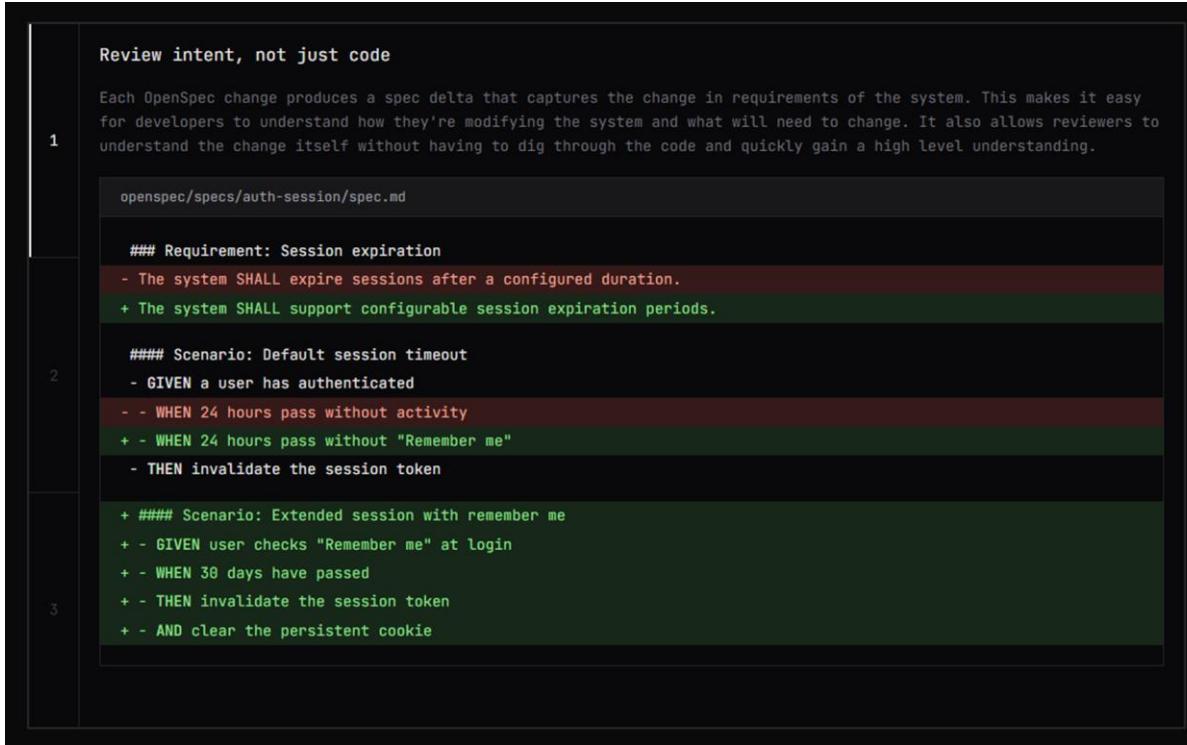
```
You: /opsx:apply
```

```
AI: Implementing tasks...  
✓ 1.1 Add theme context provider  
✓ 1.2 Create toggle component  
✓ 2.1 Add CSS variables  
✓ 2.2 Wire up localStorage  
All tasks complete!
```

```
You: /opsx:archive
```

```
AI: Archived to openspec/changes/archive/2025-01-23-add-dark-mode/  
Specs updated. Ready for the next feature.
```

# Review changes



Review intent, not just code

Each OpenSpec change produces a spec delta that captures the change in requirements of the system. This makes it easy for developers to understand how they're modifying the system and what will need to change. It also allows reviewers to understand the change itself without having to dig through the code and quickly gain a high level understanding.

1 openspec/specs/auth-session/spec.md

```
### Requirement: Session expiration
- The system SHALL expire sessions after a configured duration.
+ The system SHALL support configurable session expiration periods.
```

2 ##### Scenario: Default session timeout
- GIVEN a user has authenticated
- - WHEN 24 hours pass without activity
+ - WHEN 24 hours pass without "Remember me"
- THEN invalidate the session token

3 + ##### Scenario: Extended session with remember me
+ - GIVEN user checks "Remember me" at login
+ - WHEN 30 days have passed
+ - THEN invalidate the session token
+ - AND clear the persistent cookie



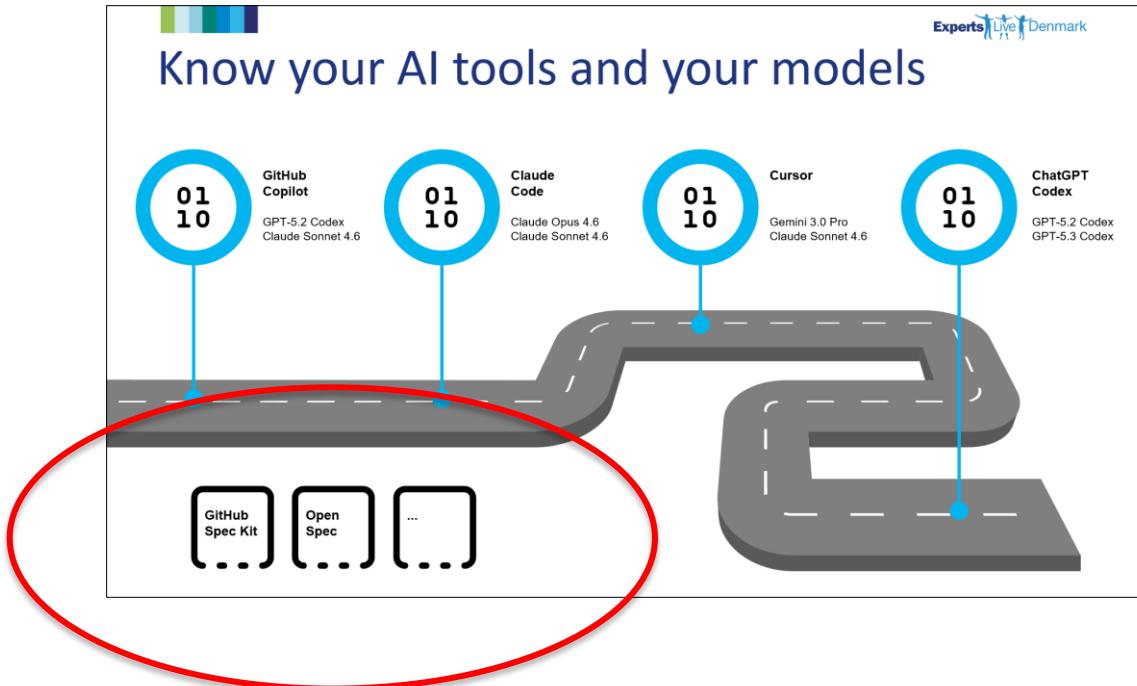
D E M O

# Spec-driven development

OpenSpec: change – implement – archive



# Spec Engines



# The change into spec-driven development

Requirements → Code ➤ Specification → Code → Validation

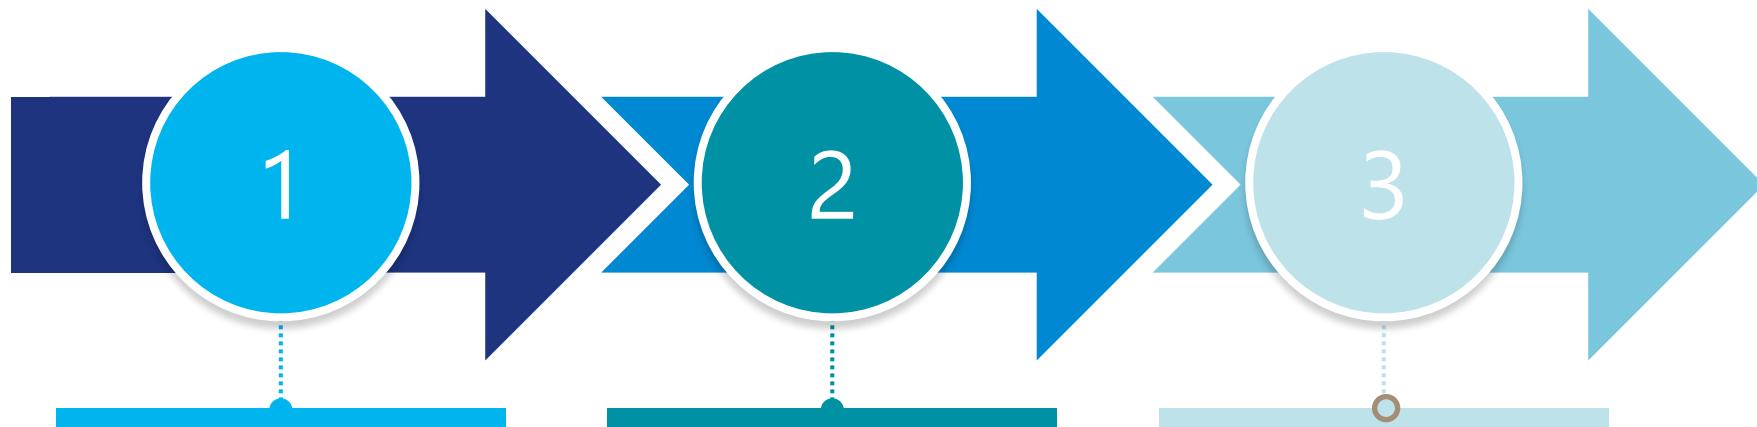
Documentation optional ➤ Specification is required

Tests define behavior ➤ Specification defines purpose and intent

AI guesses intent ➤ AI follows intent

# Agent-Driven Development

# Agent Driven Development



## KNOWLEDGE

- Description of customer, industry, and previous projects
- Description of infrastructure, governance, and compliance
- Description of project, audience, roles, and security
- Description of tech stack, architecture, and AI capabilities

## PRIMITIVES

- Import of customer, Twoday and third-party primitives
- Generation of project specific primitives including agents, instructions, chat modes, and prompts
- Compilation of primitives into standard primitives format

## ORCHESTRATION

- Use spec engine to work on features and tasks
- Follow spec-driven principals: Change => Implement => Archive => Close
- Use human orchestration and review, and AI test and review



D E M O

# Agent-driven development

Putting the pieces together



CASE

# Spec-driven development

Putting the pieces together



# Spec-Driven Development

## AI Tools and Models

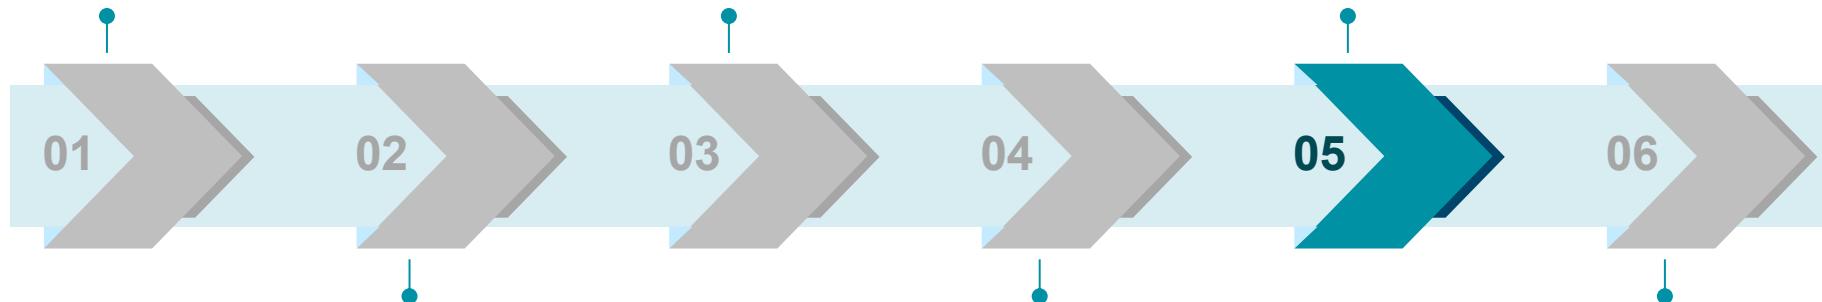
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

## AI-Assisted Development

- Prompt engineering beyond "do X"

## Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



### Vibe Coding

- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

### Agentic Workflows

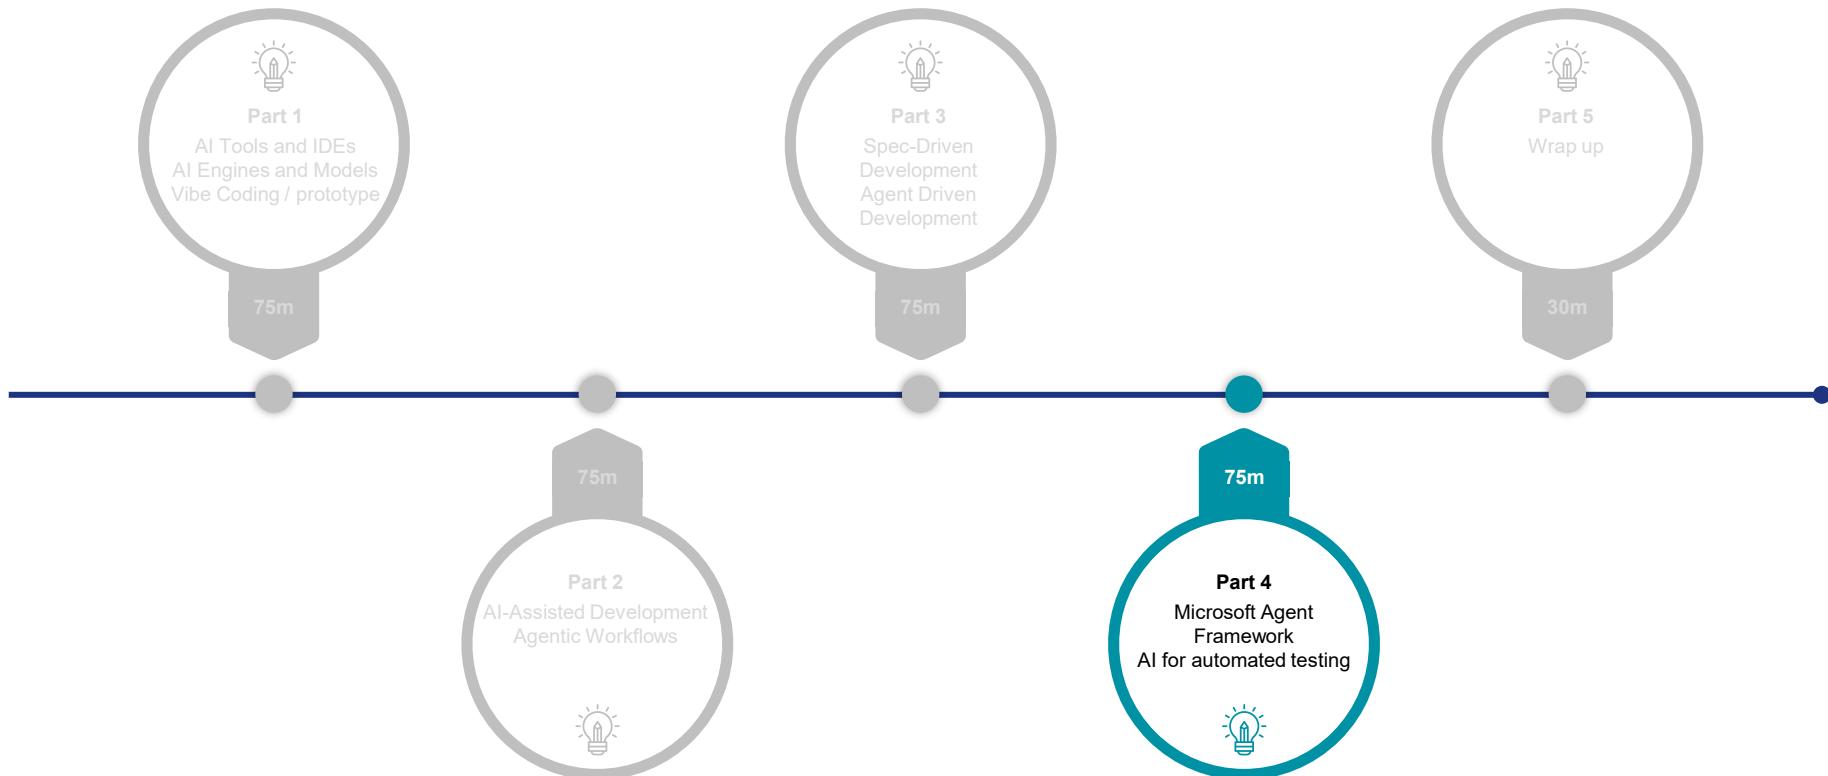
- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

### Agent Framework

- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior



# Agenda



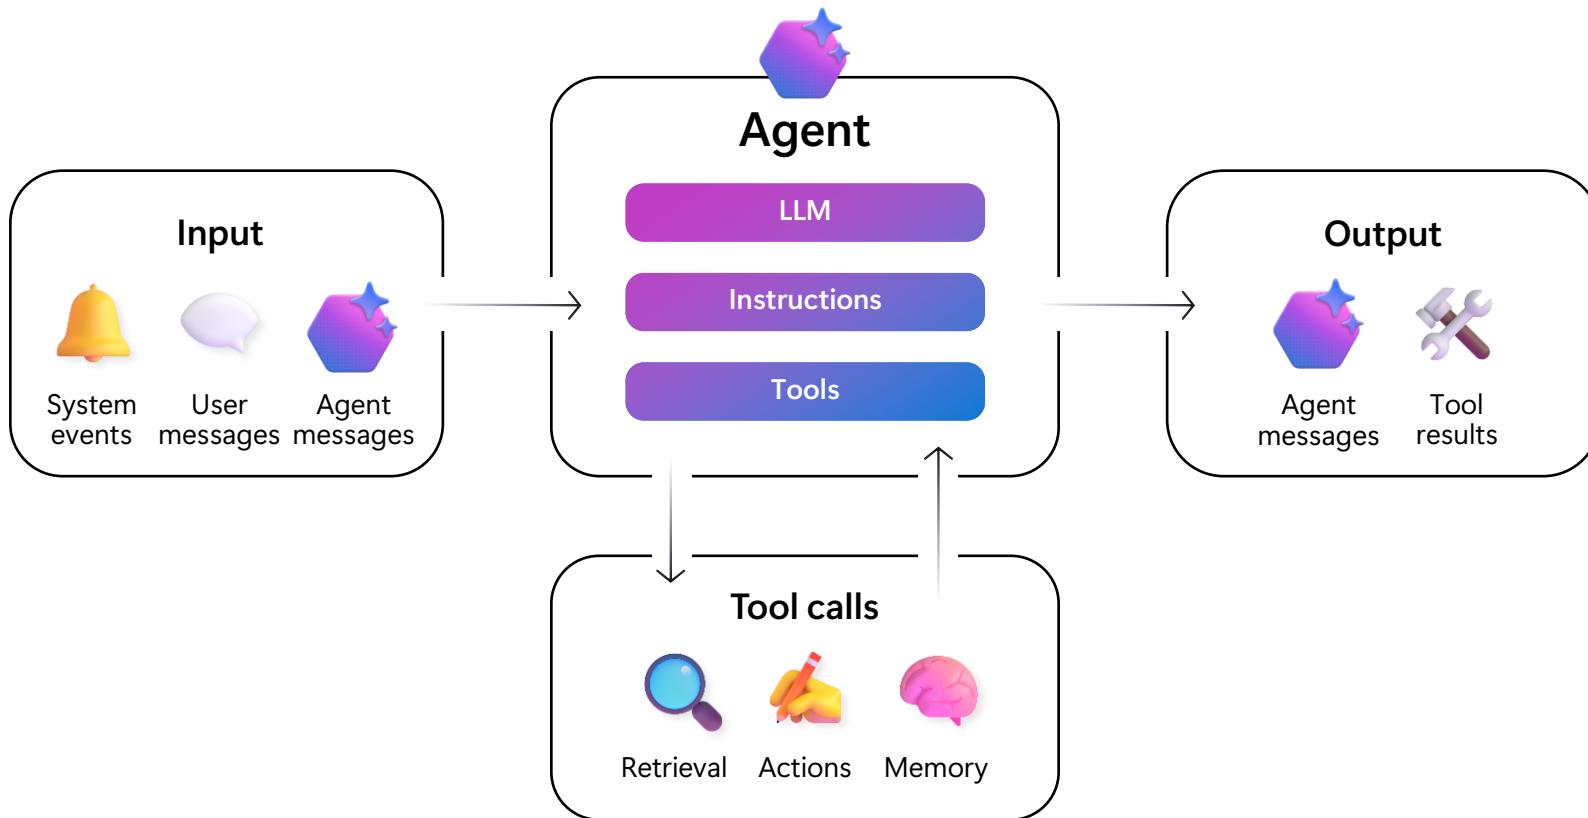


## Part 4

# Developing your own agents for AI-assisted development



# What is an agent?





## Part 4 Agenda

In the last part of this master class, we're moving on from using existing AI tools to building our own!

Plan for the rest of this part:

- Introduction of the tools and frameworks (15 min)
- Hands-on exercise (60 min)



# To make an AI agent for software development, you'll need...

Serves 1 | [Convert to metric](#)

- 1 Large Language Model for the base
- 1 handful of tools (make your own or use MCPs)
- A pinch of prompt engineering (don't skip this)
- A few thousand hours of development time to tie it all together (can be substituted with an existing framework)



Public Preview



# Microsoft Agent Framework

Open-source engine for building and orchestrating intelligent AI agents

Open Standards &  
Interoperability

Pipeline for  
Research

Community-Driven &  
Extensible by Design

Ready for  
Production

[aka.ms/AgentFramework](http://aka.ms/AgentFramework)

# Converging the two popular OSS frameworks

## Microsoft Agent Framework



### Semantic Kernel

Full SDK designed to build AI agents with ease, excellent for single agents and can be extended for multi-agents with integrations to AutoGen



### AutoGen

Powerful multi-agent research framework with pre-built conversation orchestration patterns for handling complex agent systems



# Run Agent

# Unified SDK & Simple Agent Primitives

Agent Framework introduces unified primitives (AIAgent, AgentThread, AgentTool) built on Microsoft.Extensions.AI.

Developers can create and run agents with minimal code and reuse the same abstractions across OpenAI, Foundry SDK, or any LLM backend.

## Agents

```
AIAgent agent = new AzureOpenAIClient(  
    new Uri(Environment.GetEnvironmentVariable("AZURE_OPENAI_ENDPOINT")!),  
    new AzureCliCredential()  
        .GetChatClient("gpt-4o-mini")  
        .CreateAIAgent(  
            instructions: "You are a helpful assistant, you can help the user with weather information.");  
  
Console.WriteLine(await agent.RunAsync("What's the weather in Amsterdam?"));
```

## Tools

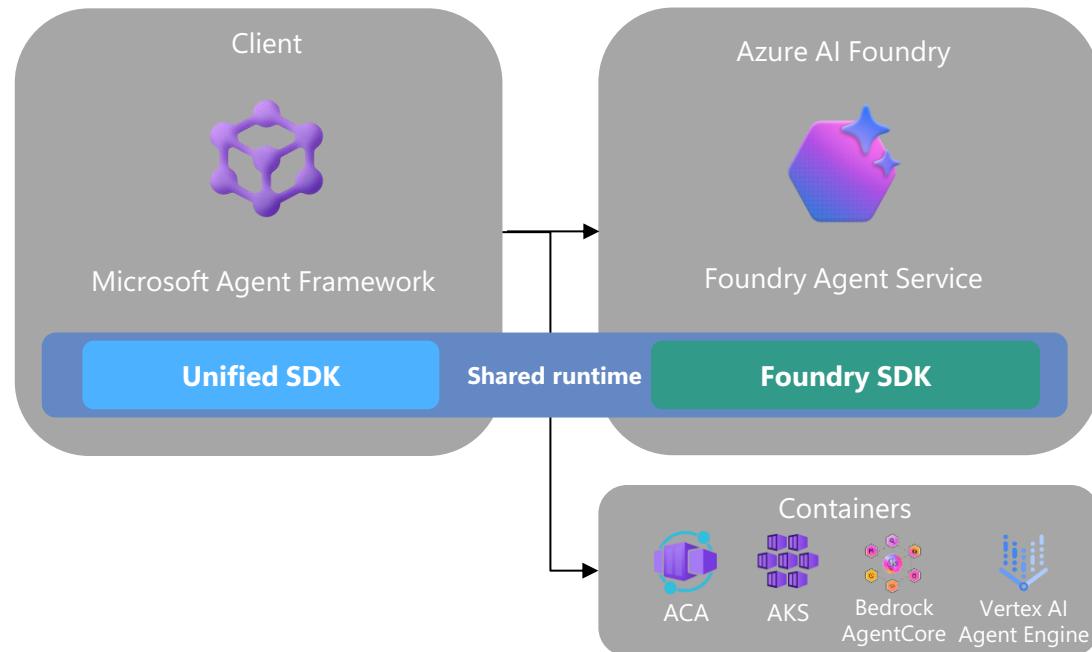
```
[Description("Get the weather for a given location.")]  
static string GetWeather([Description("The location to get the weather for.")] string location)  
=> $"The weather in {location} is cloudy with a high of 15°C."  
  
AIAgent agent = new AzureOpenAIClient(  
    new Uri(Environment.GetEnvironmentVariable("AZURE_OPENAI_ENDPOINT")!),  
    new AzureCliCredential()  
        .GetChatClient("gpt-4o-mini")  
        .CreateAIAgent(  
            instructions: "You are a helpful assistant, you can help the user with weather information.",  
            tools: [AIFunctionFactory.Create(GetWeather)]);  
  
Console.WriteLine(await agent.RunAsync("What's the weather in Amsterdam?"));
```

## Threads

```
[Description("Get the weather for a given location.")]  
static string GetWeather([Description("The location to get the weather for.")] string location)  
=> $"The weather in {location} is cloudy with a high of 15°C."  
  
AIAgent agent = new AzureOpenAIClient(  
    new Uri(Environment.GetEnvironmentVariable("AZURE_OPENAI_ENDPOINT")!),  
    new AzureCliCredential()  
        .GetChatClient("gpt-4o-mini")  
        .CreateAIAgent(  
            instructions: "You are a helpful assistant, you can help the user with weather information.",  
            tools: [AIFunctionFactory.Create(GetWeather)]);  
  
AgentThread thread = agent.GetNewThread();  
  
Console.WriteLine(await agent.RunAsync("What's the weather in Amsterdam?", thread));  
  
Console.WriteLine(await agent.RunAsync("Is it likely to rain?", thread));
```

# Local-First, Cloud-Agnostic

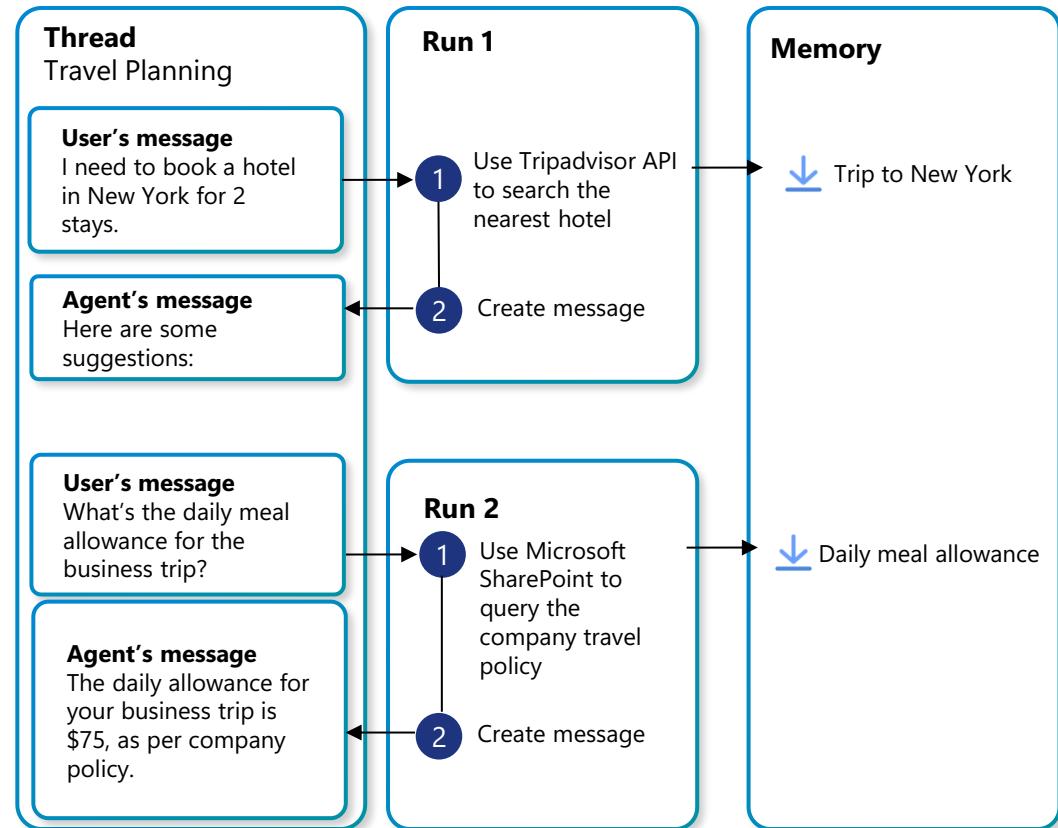
Developers can run agents locally on their laptop, test against their preferred SDK (Azure OpenAI, OpenAI SDK, Foundry SDK), and then move the same code seamlessly to Foundry Agent Service or any cloud containers without rewrites.



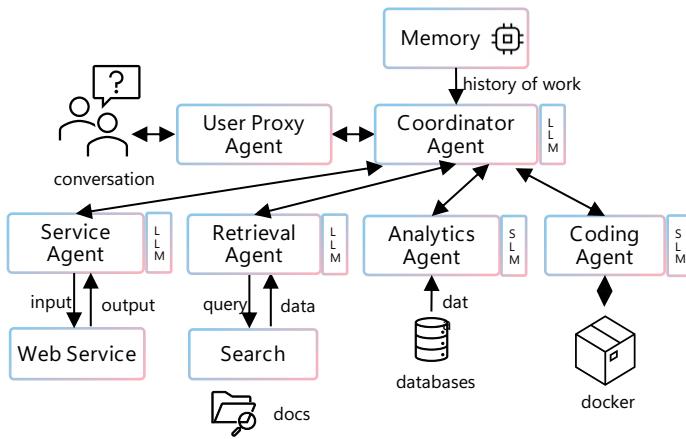
# Memory & Conversation State

The AgentThread abstraction retains conversation history across turns and sessions, ensuring agents have relevant context for long-running dialogues.

Agent Framework provides **short-term memory** (session-scoped, immediate context) and **long-term memory** (persistent data for user preferences, facts, etc.) through integrations with vector DBs.



# Agent systems with Agent Framework

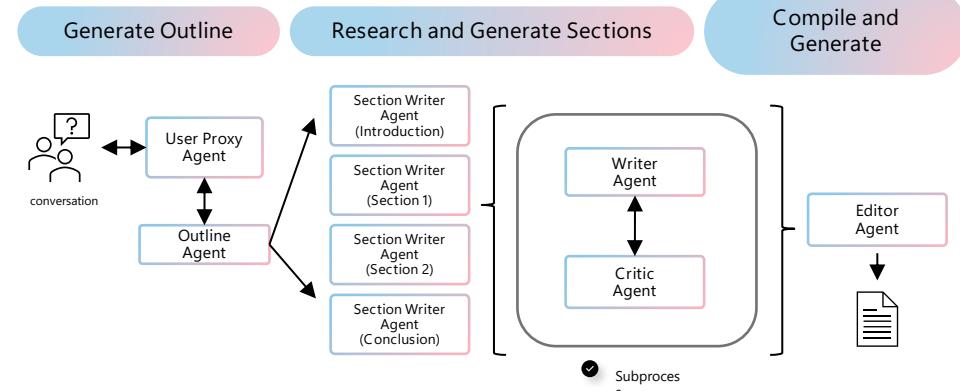


## Agent Orchestration

**LLM**-driven orchestration

*Creative reasoning and decision-making*

*Ex: Propose 2 Instagram marketing campaigns including assets that would leverage the top 2 recent trends in our past quarter US Sales to boost our mailing list user base and predict the impact of each campaign*



## Workflow Orchestration

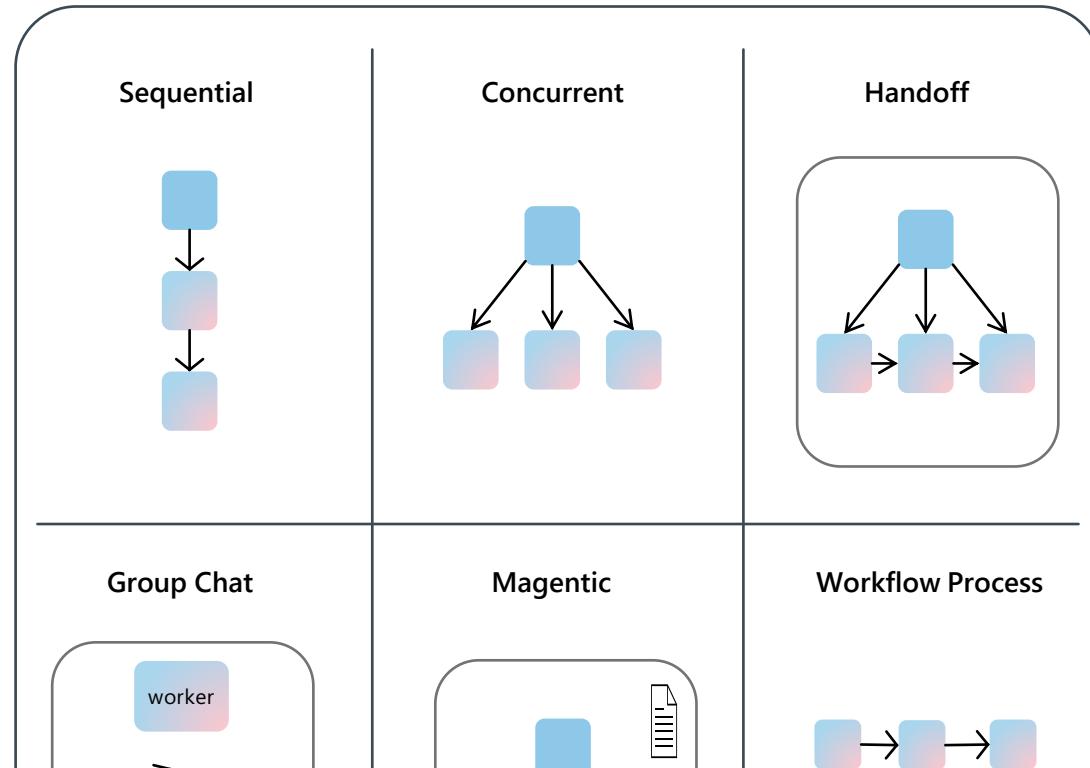
**Business logic**-driven orchestration

*Mix of business processes and agents*

*Ex: Execution the Deep Research process to generate a document outline, write and refine section and synthesize into final document*

# Multi-Agent Orchestration

Inherit advanced orchestration patterns from Microsoft Research's AutoGen combined with Semantic Kernel's durable workflow orchestration. Prototype experimental patterns locally, then scale them confidently in production.



# Tools and Extensibility

Functions, APIs, and MCP servers can all be turned into tools.

Out-of-the-box integrations to common enterprise systems, databases, and SaaS APIs — so developers don't start from scratch.

Developers can declaratively define agents in YAML and specify which tools require human approval.

## Open Ecosystem



A2A



MCP



OpenAPI



VertexAI



ElasticSearch



Pinecone



Amazon Bedrock



CrewAI



Logic Apps



Azure Functions



LangChain



MongoDB



Cosmos DB



Chroma



Qdrant



Redis



SQL Server



Faiss



Weaviate



# Introduction to Playwright MCP

[microsoft/playwright-mcp: Playwright MCP server](https://github.com/microsoft/playwright-mcp)

**Playwright enables reliable end-to-end testing for modern web apps.**

[GET STARTED](#) Star 82k+



Any browser • Any platform • One API

**Cross-browser.** Playwright supports all modern rendering engines including Chromium, WebKit, and Firefox.

**Cross-platform.** Test on Windows, Linux, and macOS, locally or on CI, headless or headed.

**Cross-language.** Use the Playwright API in [TypeScript](#), [JavaScript](#), [Python](#), [.NET](#), [Java](#).

**Test Mobile Web.** Native mobile emulation of Google Chrome for Android and Mobile Safari. The same rendering engine works on your Desktop and in the Cloud.



# Hands-on workshop

- Eight modules with Python samples on building agents
- Use Azure OpenAI authentication or ask Copilot to modify authentication for OpenAI, Anthropic, Ollama, etc.
- Finally, build an agent and test the app you developed earlier
- To get started, follow the guide at

[aka.ms/AgenticAiHack](https://aka.ms/AgenticAiHack)





# Agent Framework

## AI Tools and Models

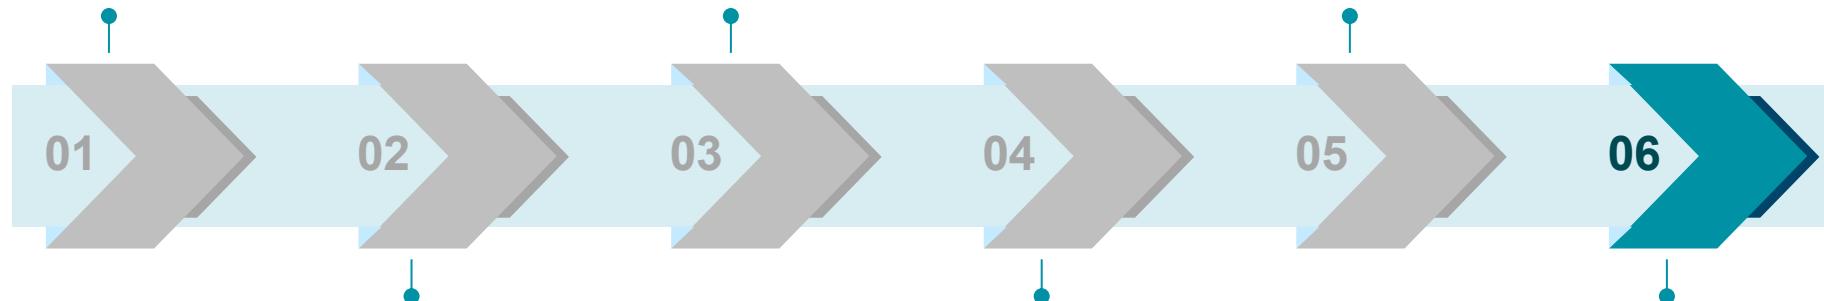
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

## AI-Assisted Development

- Prompt engineering beyond "do X"

## Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



### Vibe Coding

- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

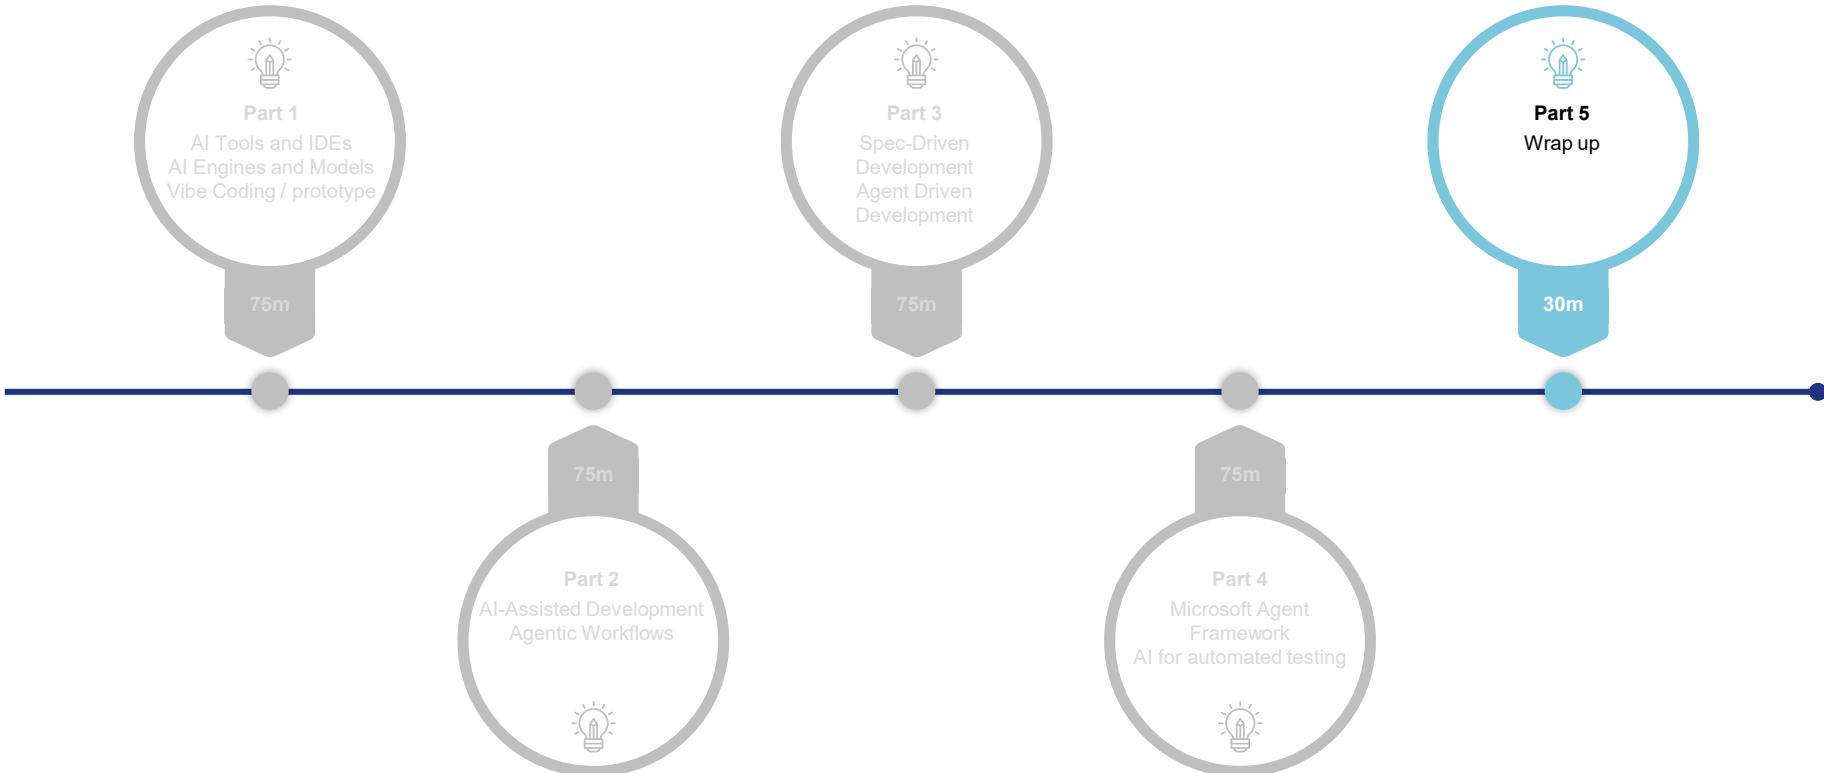
### Agentic Workflows

- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

### Agent Framework

- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior

# Agenda



# Takeaways for this master class

## Engineering shifts from coding to intent and specs

Modern AI-native development moves the engineer's focus from writing code to defining intent, specifications, architecture boundaries, and validation gates. The spec becomes the core unit of work, enabling predictable and controlled agent execution.

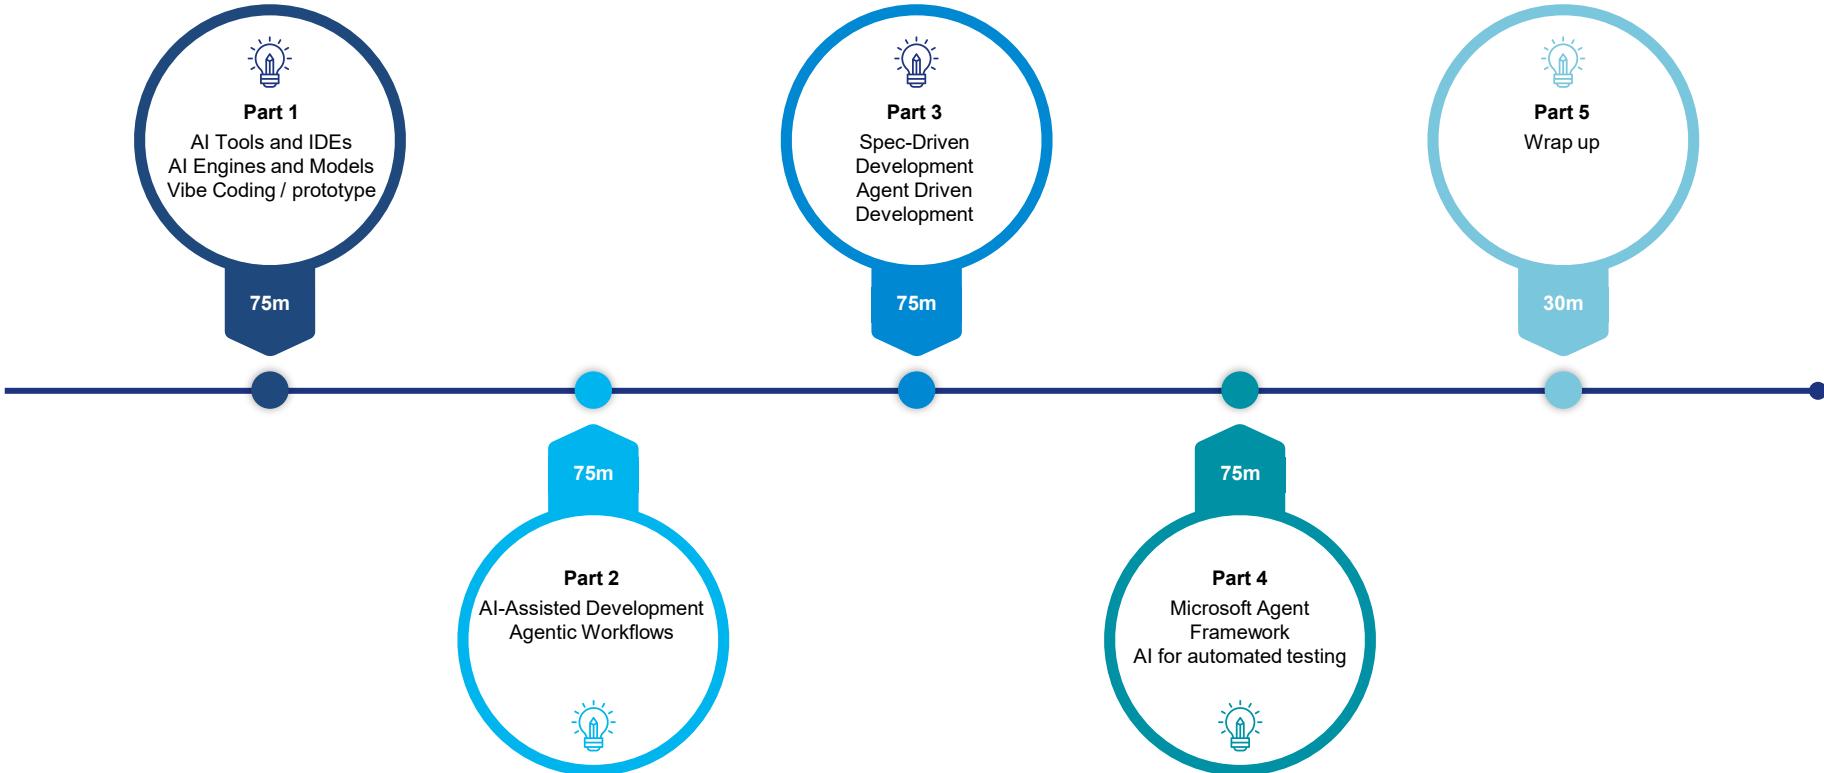
## Reliable AI development requires structure, not ad-hoc prompts

Professional AI-assisted development depends on structured prompt engineering, agent primitives, and context engineering. Treating prompts and agent behavior as reusable, versioned engineering assets makes outcomes predictable and scalable.

## AI supports the full software lifecycle—not just code

AI can do more than generate code: it can drive automated testing, documentation, validation, and other lifecycle tasks. These capabilities can be orchestrated end-to-end using interoperable frameworks like the Microsoft Agent Framework, with human oversight and validation gates.

# Agenda





## Treadley Case Overview

### AI Tools and Models

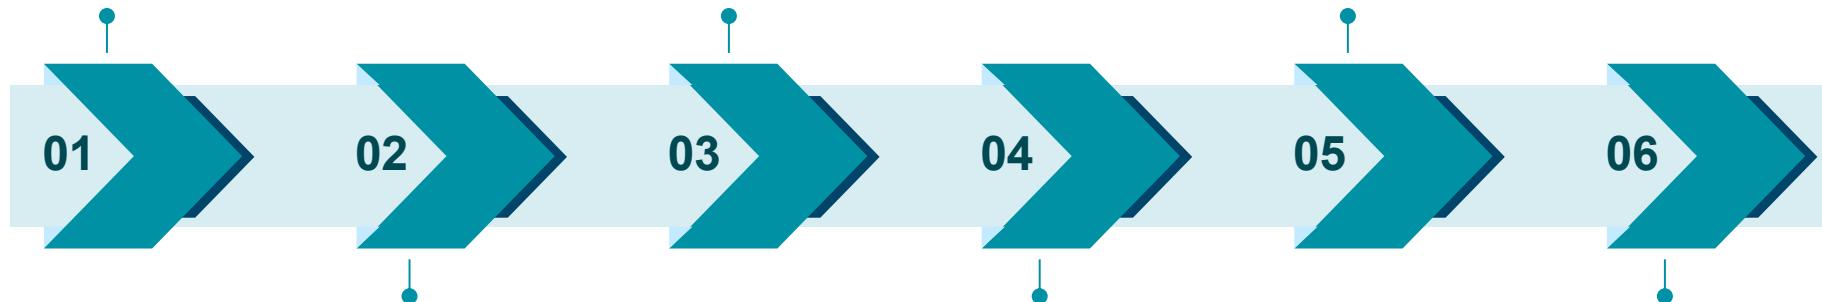
- Using AI tools and IDE integrations effectively
- Understanding model strengths and limitations

### AI-Assisted Development

- Prompt engineering beyond "do X"

### Spec-Driven Development

- Writing specs that are AI-readable and human-readable
- Separating intent from implementation
- Using specs as a stable contract



### Vibe Coding

- Co-creating code with AI, not over-planning
- Letting AI "fill in the blanks" responsibly

### Agentic Workflows

- Agent primitives (plan, act, observe, reflect)
- Context engineering (what the AI needs to know)
- Multi-step, goal-oriented workflows

### Agent Framework

- Agent framework concepts
- Tooling and orchestration
- Testing AI-driven systems
- Validating code and agent behavior



DEMO

T-shirt shops





# Session Feedback – THANK YOU 😊

Please evaluate this session using QR Code **during next 30 min**

- or -

Make your session feedback on your way out of the room



<https://smiley.link/CFCJPFCG>